Principles of Database Systems (CS307)

Yuxin Ma

Department of Computer Science and Engineering
Southern University of Science and Technology

Instructors

- **Yuxin Ma** 马昱欣 (Associate Professor in the Department of CSE)
  - Office: Room 512, South Tower, Engineering Building
  - Personal Website
    - https://cse.sustech.edu.cn/faculty/~mayx/
    - https://mayuxin.me (mirror site)
  - Email: mayx@sustech.edu.cn

- **Yueming Zhu** 朱悦铭 (Lab Lecturer in the Department of CSE)
  - Office: Room 111, South Tower, Engineering Building
  - Email: zhuym@sustech.edu.cn

Office Hour

- **Yuxin Ma**: Monday 14:00 – 16:00 @ Office
- **Yueming Zhu**: Wednesday 9:30 – 11:30 @ Office

Lecture Notes, Blackboard and WeCom Group

- There is a Blackboard site for this course
  - Lecture slides, lab sheets, assignments, and project requirements will be released on the Blackboard site

- Lab sessions
  - Tuesday 7-8 (Yuxin Ma)
  - Tuesday 7-8 (Yueming Zhu)
  - Wednesday 5-6 (Yueming Zhu)

Lecture Notes, Blackboard and WeCom Group

- QR code for the SUSTech WeCom (企业微信) group
  - Teaching assistants will be there to help you

Textbooks

- Reference book:
  - A. Silberschatz, H. Korth, and S. Sudarshan. **Database System Concepts**. McGraw-Hill, New York, 7th Edition, (2019).

Grading Policy

- Lecture and Lab Attendance (10%)
  - 5% for each part
- Assignments (15%)
- Project (35%)
  - 2 Projects
- Final exam (40%)

Principles of Database Systems (CS307)

Lecture 1: Introduction to Databases

Yuxin Ma

Department of Computer Science and Engineering
Southern University of Science and Technology

- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7th Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.

What is a Database?

- Well, first, let's take a step back: What is data?

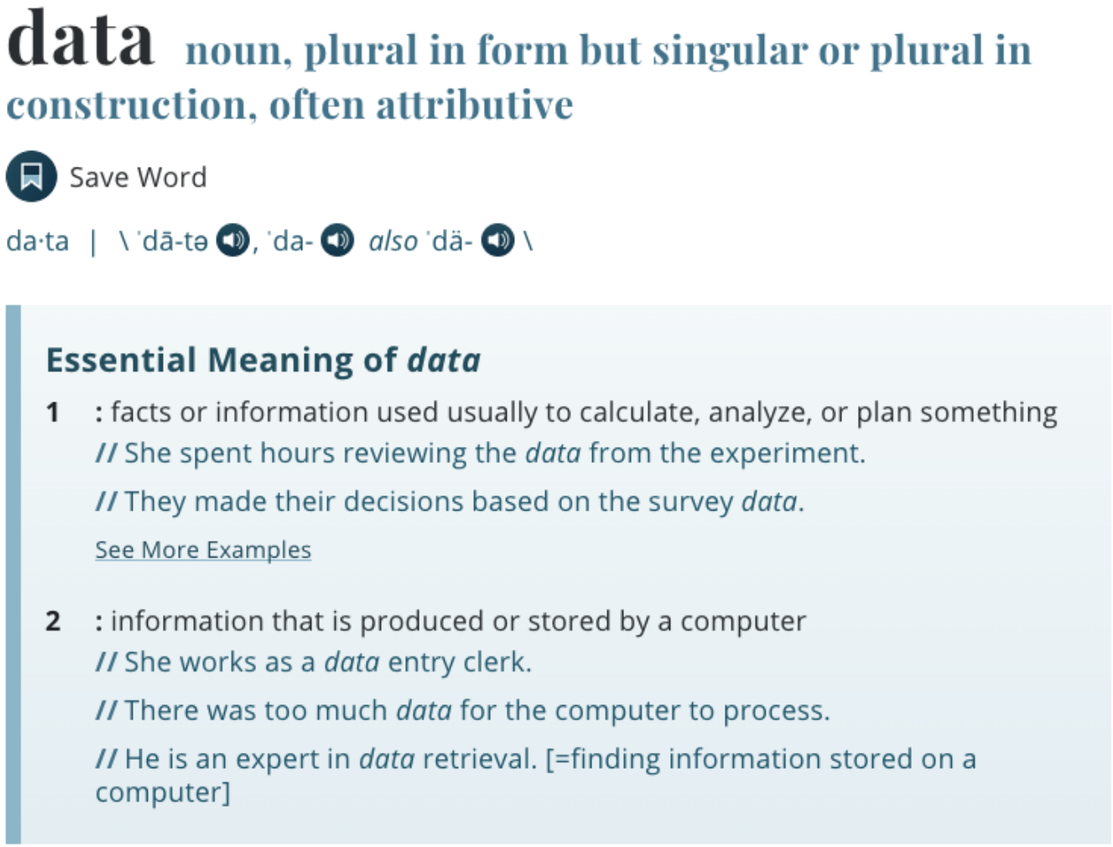

data noun, plural in form but singular or plural in construction, often attributive

Save Word

da·ta | \ ˈdā-tə, ˈda- *also* ˈdä- \

Essential Meaning of *data*

1 : facts or information used usually to calculate, analyze, or plan something
// She spent hours reviewing the *data* from the experiment.
// They made their decisions based on the survey *data*.
See More Examples

2 : information that is produced or stored by a computer
// She works as a *data* entry clerk.
// There was too much *data* for the computer to process.
// He is an expert in *data* retrieval. [=finding information stored on a computer]

What is a Database?

- A modern database system is a complex software system whose task is to manage a large, complex collection of data.
  - Collection of interrelated data
  - Set of programs to access the data
  - An environment that is both *convenient* and *efficient* to use
- Databases touch all aspects of our lives

Applications of Database

- Enterprise Information
  - **Sales**: customers, products, purchases
  - **Accounting**: payments, receipts, assets
  - **Human Resources**: Information about employees, salaries, payroll taxes.
- Manufacturing
  - Management of production, inventory, orders, supply chain.
- Universities
  - Registration, grades

Applications of Database

- Databases are everywhere today
  - … but the concept is old
  - The idea was to have one system doing once and for all the boring data storage/retrieval part
    - What boring parts?

Purpose of Database Systems

- In the early days …
  - Database applications were built directly on top of file systems
  - However, it suffers from many issues, including (but not limited to):
    - Data redundancy and inconsistency
    - Difficulty in accessing data
    - Data isolation
    - Integrity problems
    - Atomicity of updates
    - Concurrent access by multiple users
    - Security problems

Let’s Show an Example

- Write a Java program to manage information of all students in CS307
  - We have classroom sessions and lab sessions
    - (potential redundancy of students)
  - A new Java class method for each task
    - (difficulty in accessing data)
  - Maybe you will split the students into files based on lab session times
    - (split into multiple files; hard to manage the files)
  - You need to check the validity of Student IDs and classes
    - (hard to maintain integrity constraints)
  - Atomicity? Concurrency? Security? All need to be handled by yourself

Let's Show an Example (Well, still remember how to read and write files in Java?)

- Write a Java program to manage information of all students in CS307
  - We have classroom sessions and lab sessions
    - (potential redundancy of students)
  - A new Java class method for each task
    - (difficulty in accessing data)
  - Maybe you will split the students into files based on lab session times
    - (split into multiple files; hard to manage the files)
  - You need to check the validity of Student IDs and classes
    - (hard to maintain integrity constraints)
  - Atomicity? Concurrency? Security? All need to be handled by yourself

Let’s Show an Example

- Furthermore,
  - Can you reuse the code you just wrote in a staff management system?

Purpose of Database Systems

- In the early days …
  - Database applications were built directly on top of file systems

  - However, it suffers from many issues, including (but not limited to):
    - Data redundancy and inconsistency
    - Difficulty in accessing data
    - Data isolation
    - Integrity problems
    - Atomicity of updates
    - Concurrent access by multiple users
    - Security problems

Database systems offer solutions to all the problems mentioned above

A Bit of History

- 1950s and early 1960s:
  - Data processing using magnetic tapes for storage
    - Tapes provided only sequential access
  - Punched cards for input
- Late 1960s and 1970s:
  - Hard disks allowed direct access to data
  - Network and hierarchical data models in widespread use
  - **Ted Codd** defines the relational data model
    - Would win the ACM Turing Award for this work
    - IBM Research begins System R prototype
    - UC Berkeley (Michael Stonebraker) begins Ingres prototype
    - Oracle releases first commercial relational database
  - High-performance (for the era) transaction processing

A Bit of History

Edgar F. “Ted” Codd
1923 – 2003
Turing Award 1981

A Relational Model of Data for Large Shared Data Banks

E. F. Codd
IBM Research Laboratory, San Jose, California

Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation). A prompting service which supplies such information is not a satisfactory solution. Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed. Changes in data representation will often be needed as a result of changes in query, update, and report traffic and natural growth in the types of stored information.

Existing noninferential, formatted data systems provide users with tree-structured files or slightly more general network models of the data. In Section 1, inadequacies of these models are discussed. A model based on *n*-ary relations, a normal form for data base relations, and the concept of a universal

The relational view (or model) of data described in Section 1 appears to be superior in several respects to the graph or network model [3, 4] presently in vogue for non-inferential systems. It provides a means of describing data with its natural structure only—that is, without superimposing any additional structure for machine representation purposes. Accordingly, it provides a basis for a high level data language which will yield maximal independence between programs on the one hand and machine representation and organization of data on the other.

A further advantage of the relational view is that it forms a sound basis for treating derivability, redundancy, and consistency of relations—these are discussed in Section 2. The network model, on the other hand, has spawned a number of confusions, not the least of which is mistaking the derivation of connections for the derivation of relations (see remarks in Section 2 on the “connection trap”).

Finally, the relational view permits a clearer evaluation of the scope and logical limitations of present formatted data systems, and also the relative merits (from a logical standpoint) of competing representations of data within a single system. Examples of this clearer perspective are cited in various parts of this paper. Implementations of systems to support the relational model are not discussed.

E. F. Codd, A Relational Model of Data for Large Shared Data Banks, Information Retrieval, June, 1970

A Bit of History

- 1980s:
  - Research relational prototypes evolve into commercial systems
    - SQL becomes industrial standard
  - Parallel and distributed database systems
    - Wisconsin, IBM, Teradata
  - Object-oriented database systems
- 1990s:
  - Large decision support and data-mining applications
  - Large multi-terabyte data warehouses
  - Emergence of Web commerce

A Bit of History

- 2000s
  - Big data storage systems
    - Google BigTable, Yahoo PNuts, Amazon,
    - “NoSQL” systems.
  - Big data analysis: beyond SQL
    - Map reduce and friends
- 2010s
  - SQL reloaded
    - SQL front end to Map Reduce systems
    - Massively parallel database systems
    - Multi-core in-memory databases

Relational Database

- Based on the relational model of data
  - Organizes data into one or more tables
  - Rows are also called records or tuples
  - Columns are also called attributes

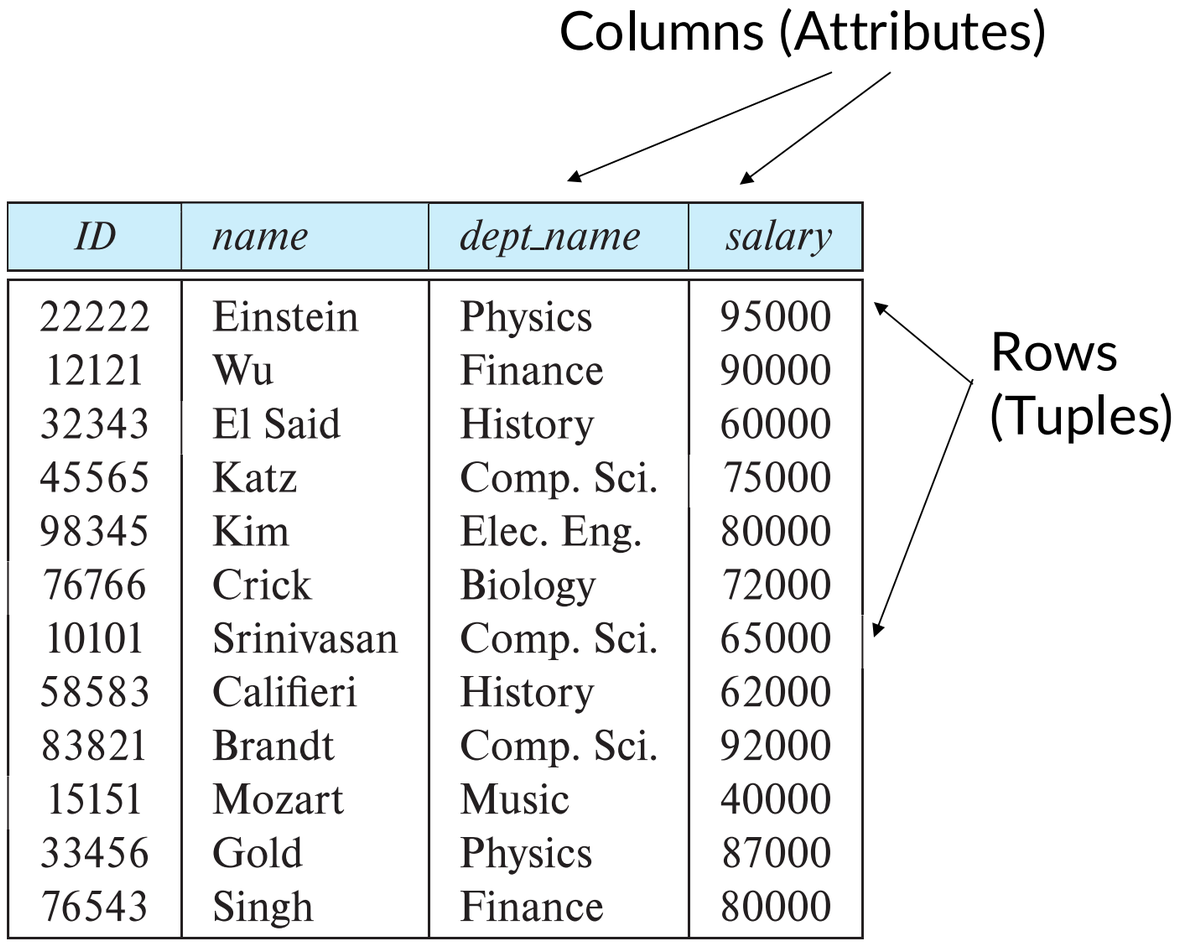

ID	*name*	*dept_name*	*salary*
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

Relational Database

- Each column stores a piece of data

- One row represents a “known fact”:
  - “Audrey Hepburn was born on 1929/05/04 and looked like this.”

Surname	Firstname	Birthdate	Picture
Hepburn	Audrey	4-May-1929	

Relational Database

- Each column stores a piece of data
- One row represents a “known fact”:
  - “Audrey Hepburn was born on 1929/05/04 and looked like this.”
- All the pieces of data in a row are related, hence "**relational**"

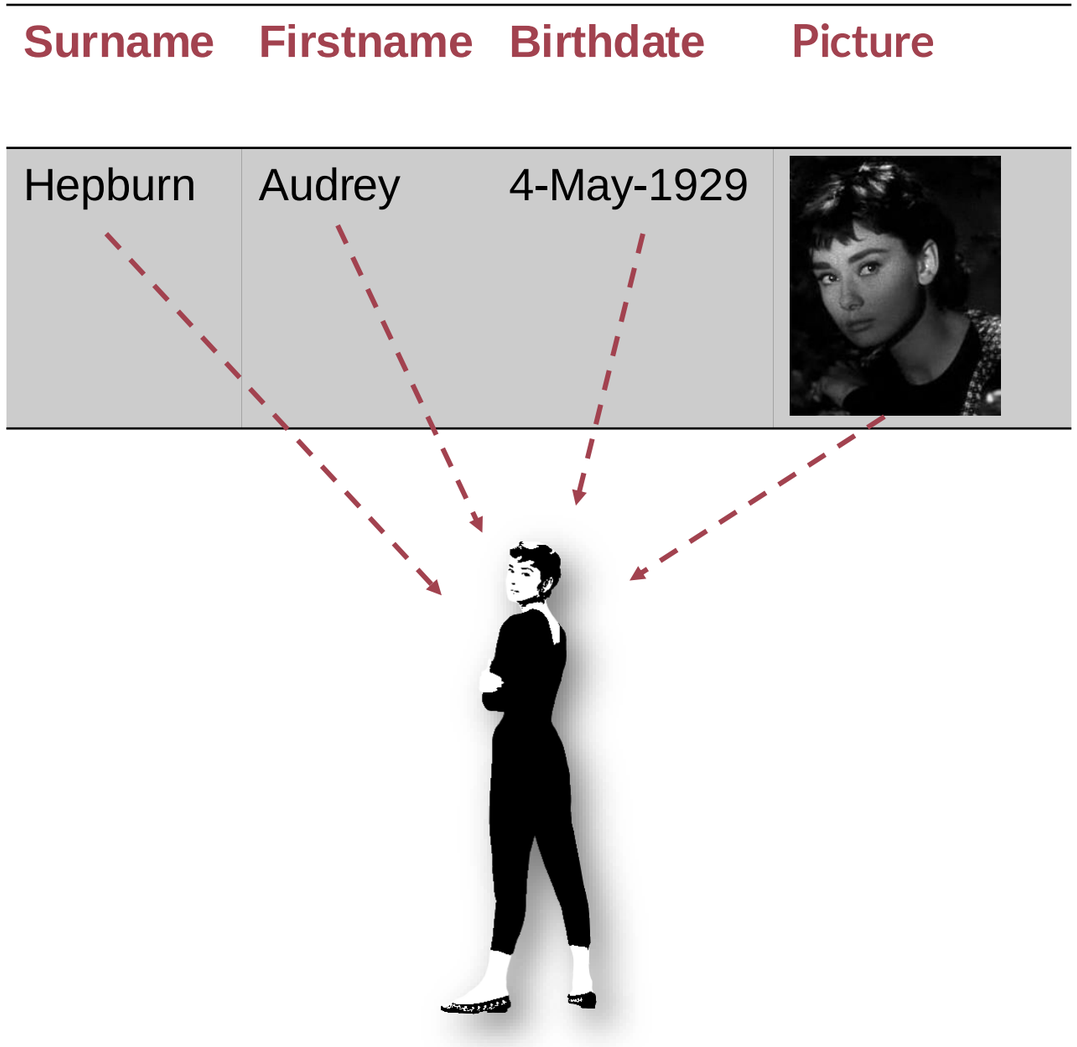

Relational Database

- But Codd's "big idea"
  - You could operate on the relations and get new sets
- Relational Algebra
  - Theoretical foundation for relational databases

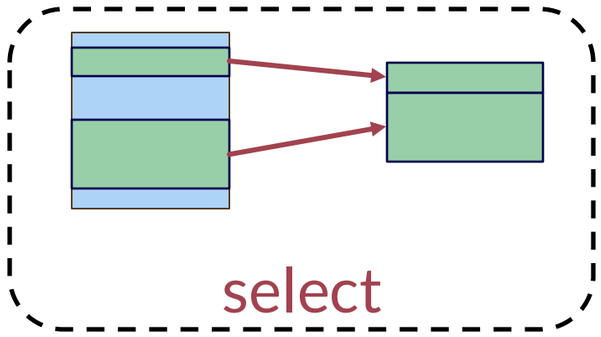

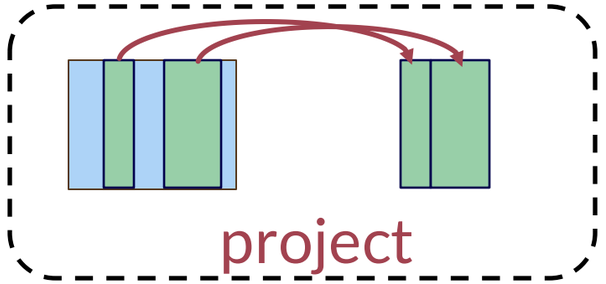

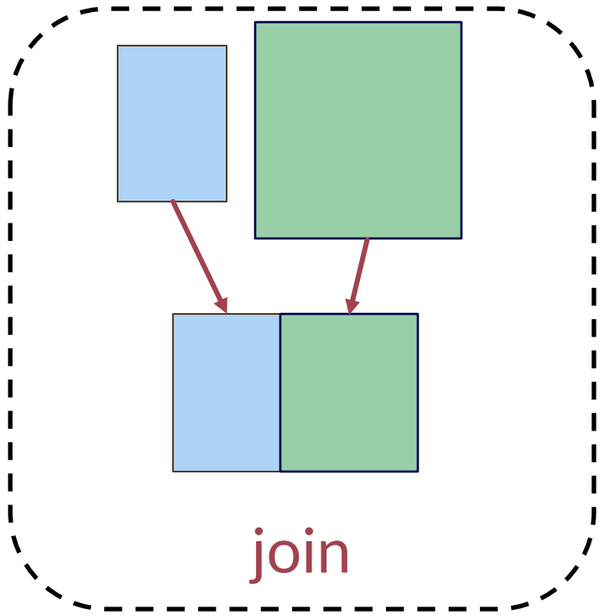

Key

- Example: A Film Database
  - Easy to find such as “The 100 greatest films ever”
  - Sometimes as a `.csv` file that you can load into a spreadsheet

1	Movie Title	Country	Year	Director	Starring
2	Citizen Kane	US	1941	welles, o.	Orson Welles, joseph Cotten
3	La règle du jeu	FR	1939	Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
4	North by Northwest	US	1959	HITCHCOCK, A.	Cary GRANT, Eva Marie SAINT, James MASON
5	Singin' In the Rain	US	1952	Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
6	Rear Window	US	1954	HITCHCOCK, A.	James STEWART, Grace KELLY
7	City Lights	US	1931	CHAPLIN, C.	Charlie CHAPLIN, Virginia CHERRILL
8	The Third Man	GB	1949	Reed, C.	Joseph Cotten, Alida Valli, Orson Welles
9	The Searchers	US	1956	Ford, J.	John Wayne, Jeffrey Hunter, Natalie Wood
10	Ladri di biciclette	IT	1949	DeSica, V.	Lamberto Maggiorani, Enzo Staiola
11	Annie Hall	US	1977	Allen, W.	Woody Allen, Diane Keaton
12	On the Waterfront	US	1954	Kazan, E.	Marlon Brando, Eva Marie Saint, Karl Malden
13	All about Eve	US	1950	Mankiewicz, J.	Bette Davis, Anne Baxter, George Sanders
14	Casablanca	US	1942	Curtiz, M.	Humphrey Bogart,Ingrid Bergman, Claude Rains
15	The Treasure of the Sierra Madre	US	1948	HUSTON, J.	Humphrey BOGART, Walter HUSTON, Tim HOLT
16	High Noon	US	1952	Zinnemann, F.	Gary Cooper, Grace Kelly
17	Some Like It Hot	US	1959	Wilder, B.	Tony Curtis, Jack Lemmon, Marilyn Monroe
18					

Key

- Duplicates are forbidden in relational tables
  - Or, it would introduce potential errors, such as in counting the number of movies

Movie Title	Country	Year	Director	Starring
Citizen Kane	US	1941	welles, o.	Orson Welles, Joseph Cotten
La règle du jeu	FR	1939	Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
North By Northwest	US	1959	HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
Singin' in the Rain	US	1952	Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
Rear Window	US	1954	HITCHCOCK, A.	James Stewart, Grace Kelly
North By Northwest	US	1959	HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason

Key

- How to identify “different rows” in a table?
  - A column (or a set of columns) to differentiate one row from another
    - In the film data … How about the movie titles?

Movie Title	Country	Year	Director	Starring
Citizen Kane	US	1941	welles, o.	Orson Welles, Joseph Cotten
La règle du jeu	FR	1939	Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
North By Northwest	US	1959	HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
Singin' in the Rain	US	1952	Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
Rear Window	US	1954	HITCHCOCK, A.	James Stewart, Grace Kelly

豆瓣电影 神雕侠侣

影讯&购票 选电影 电视剧 排行榜 分类 影评 2021年度榜单 2021书影音报告

搜索 神雕侠侣

神雕侠侣 神雕侠侣 (1995) [剧集] [可播放]

★★★★½ 9.2 (176564人评价)

中国香港 / 爱情 / 武侠 / 古装 / 新神雕侠侣 / Return Of The Condor Heroes / 45分钟

李添胜 / 古天乐 / 李若彤 / 白彪 / 魏秋桦 / 傅明宪 / 李绮虹 / 雪梨 / 简佩玛

话题 《神雕侠侣》哪个角色最让你惊喜?

10030人浏览 · 22篇文章

神雕侠侣 (2006) [剧集] [可播放]

★★★★☆ 7.5 (70883人评价)

中国大陆 / 武侠 / 古装 / The Return of the Condor Heroes / 40分钟

于敏 / 刘亦菲 / 黄晓明 / 陈紫函 / 杨幂 / 叮当 / 王洛勇 / 赵鸿飞 / 钱博

神雕侠侣 (2014) [剧集] [可播放]

★★½☆☆ 4.9 (26245人评价)

中国大陆 / 剧情 / 武侠 / 新神雕侠侣 / The Condor Heroes / 45分钟

李慧珠 / 邓伟恩 / 李达超 / 陈晓 / 陈妍希 / 张馨予 / 张雪迎 / 郑国霖 / 杨明娜 / 陈翔 / 毛晓彤

新神雕侠侣 (2022) [剧集]

☆☆☆☆☆ (尚未播出)

中国大陆 / 爱情 / 武侠 / 古装

林峰 / 佟梦实 / 毛晓慧 / 文淇 / 涂冰 / 邵兵 / 龚蓓苾 / 毛林林 / 宗峰岩

神雕侠侣 (2023)

☆☆☆☆☆ (尚未上映)

中国大陆 / 剧情 / 爱情 / 武侠 / 古装 / The Romance Of The Condor Heros

徐克

神雕侠侣 神雕侠侣 (1983) [剧集] [可播放]

★★★★☆ 7.9 (6328人评价)

中国香港 / 剧情 / 动作 / 爱情 / The Return of the Condor Heroes / 42分钟

范秀明 / 鞠觉亮 / 萧显辉 / 司徒立光 / 谭锐铭 / 刘德华 / 陈玉莲 / 梁家仁 / 欧阳佩珊 / 廖安丽 / 吕有慧 / 曾江

神雕侠侣 (1998) [剧集] [可播放]

★★★½☆ 7.1 (5767人评价)

新加坡 / 爱情 / 武侠 / 古装 / 神雕侠侣 98版 / Return of the Condor Heroes / 47分钟

马玉辉 / 谢敏洋 / 蔡晶盛 / 张龙敏 / 卢燕金 / 李铭顺 / 范文芳 / 朱厚任 / 何咏芳 / 林湘萍 / 丁岚 / 李南星

神雕侠侣 (1998) [剧集] [可播放]

★★½☆☆ 5.2 (10379人评价)

中国台湾 / 中国大陆 / 武侠 / 古装 / 杨过与小龙女 / 45分钟

李惠民 / 赖水清 / 任贤齐 / 吴倩莲 / 孙兴 / 季芹 / 李立群 / 夏文汐 / 蔡君茹 / 高捷

神雕侠侣 (2001) [剧集] [可播放]

★★★★☆ 7.5 (1772人评价)

日本 / 中国香港 / 动画 / 24分钟

案纳正美 / 高木淳 / 浪川大辅 / 园崎未惠 / 中田让治 / 唐泽润 / 木村亚希子 / 小村哲生 / 广田行生 / 高户靖广

神雕侠侣 (1984) [剧集]

★★★★☆ 7.5 (926人评价)

中国台湾 / 剧情 / 爱情 / 武侠 / 古装 / 95分钟

何东兴 / 孟飞 / 潘迎紫 / 傅娟

九一神雕侠侣 九一神鵰侠侶 (1991) [可播放]

★★★☆☆ 6.4 (8789人评价)

中国香港 / 动作 / 剧情 / 奇幻 / 爱情 / 科幻 / 神秘英豪 / 新神雕侠侣 / 92分钟

元奎 / 黎大炜 / 刘德华 / 梅艳芳 / 郭富城 / 叶蕴仪 / 刘嘉玲

神雕侠侣 神鵰侠侶 (1982) [可播放]

★★½☆☆ 5.2 (1402人评价)

中国香港 / 动作 / 爱情 / 武侠 / 古装 / 射雕英雄传4 / Brave Archer 4 / 100分钟

张彻 / 江生 / 郭追 / 傅声 / 龙天翔 / 黄淑仪

九二神雕之痴心情长剑 九二神鵰之痴心情長劍 (1992) [可播放]

★★★☆☆ 6.0 (6152人评价)

中国香港 / 喜剧 / 爱情 / 奇幻 / 武侠 / 神秘情侠 / 新神雕侠侣2(台) / 92分钟

元奎 / 黎大炜 / 刘德华 / 关之琳 / 吴耀汉 / 关淑怡

神雕侠侣 (1976) [剧集]

☆☆☆☆☆ (暂无评分)

中国香港 / 古装

萧笙 / 罗乐林 / 李通明 / 白彪 / 米雪 / 曾江 Kenneth Tsang / 秦煌 / 郑裕玲 / 冯淬帆

Key

- How to identify “different rows” in a table?
  - A column (or a set of columns) to differentiate one row from another
    - In the film data … ~~How about the movie titles?~~ Title + Director?

Movie Title	Country	Year	Director	Starring
Citizen Kane	US	1941	welles, o.	Orson Welles, Joseph Cotten
La règle du jeu	FR	1939	Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
North By Northwest	US	1959	HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
Singin' in the Rain	US	1952	Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
Rear Window	US	1954	HITCHCOCK, A.	James Stewart, Grace Kelly

But, Hitchcock did It (So it was with the COD game series)

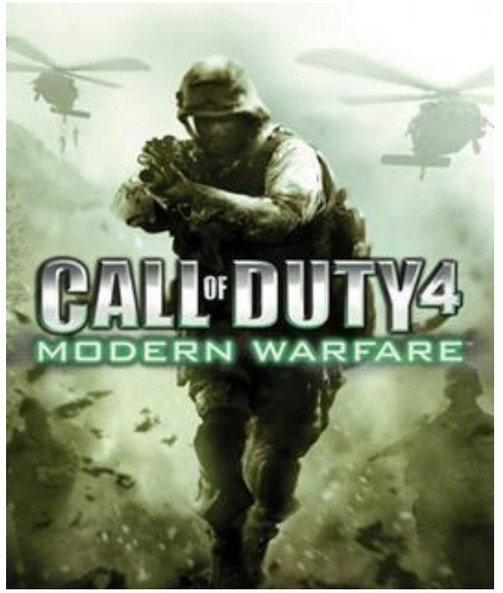

Same developer (Infinite Ward)
Same publisher (Activision)
Two different years (2007, 2019)

Key

- How to identify “different rows” in a table?
  - A column (or a set of columns) to differentiate one row from another
    - In the film data … ~~How about the movie titles? Title + Director?~~
    - Title + Director + Year?

Movie Title	Country	Year	Director	Starring
Citizen Kane	US	1941	welles, o.	Orson Welles, Joseph Cotten
La règle du jeu	FR	1939	Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
North By Northwest	US	1959	HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
Singin' in the Rain	US	1952	Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
Rear Window	US	1954	HITCHCOCK, A.	James Stewart, Grace Kelly

Key

- Ok, you have made it this time, but –
  - The combination is too difficult for either remembering or computing
    - Need to compare multiple times on different columns
    - Think about deduplication
  - What if there are multiple items in a single column?
    - For example, more than one directors in a movie?

Primary Key

Additional material about creating a unique ID for each row:
`https://en.wikipedia.org/wiki/Universally_unique_identifier`

- Some of the keys may be unique for every row
  - Student ID, Email address, 18-digit ID number, etc.
- Usually, it is a good practice to choose the simplest one
  - (Or, create one)

Movie ID	Movie Title	Country	Year	Director	Starring
0	Citizen Kane	US	1941	welles, o.	Orson Welles, Joseph Cotten
1	La règle du jeu	FR	1939	Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
2	North By Northwest	US	1959	HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
3	Singin' in the Rain	US	1952	Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
4	Rear Window	US	1954	HITCHCOCK, A.	James Stewart, Grace Kelly

Normalization

- A way of standardize your data

Movie ID	Movie Title	Country	Year	Director	Starring
0	Citizen Kane	US	1941	welles, o.	Orson Welles, Joseph Cotten
1	La règle du jeu	FR	1939	Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
2	North By Northwest	US	1959	HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
3	Singin' in the Rain	US	1952	Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
4	Rear Window	US	1954	Alfred Hitchcock	James Stewart, Grace Kelly

Too many different ways of spelling

Normalization

- “First Norm Rule” (1NF)
  - Each column should only contain ONE piece of information

Director	Starring
welles, o.	Orson Welles, Joseph Cotten
Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
HITCHCOCK, A.	James Stewart, Grace Kelly

Director_Firstname	Director_Lastname
Alfred	Hitchcock
Orson	Welles

Starring_Firstname	Starring_Lastname
Orson	Welles
Joseph	Cotten

Normalization

- “First Norm Rule” (1NF)
  - Each column should only contain ONE piece of information

Director	Starring
welles, o.	Orson Welles, Joseph Cotten
Renoir, J.	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
HITCHCOCK, A.	Cary Grant, Eva Marie Saint, James Mason
Donen/Kelly	Gene Kelly, Debbie Reynolds, Donald O'Connor
HITCHCOCK, A.	James Stewart, Grace Kelly

Director_ Firstname	Director_ Lastname	Born	Died
Alfred	Hitchcock	1899	1980
Orson	Welles	1915	1985

Extend the table to represent all directors

Normalization

- “First Norm Rule” (1NF)
  - Each column should only contain ONE piece of information

Movie ID	Movie Title	Country	Year	Director ID	Starring
0	Citizen Kane	US	1941	2	Orson Welles, Joseph Cotten
1	La règle du jeu	FR	1939	5	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
2	North By Northwest	US	1959	1	Cary Grant, Eva Marie Saint, James Mason
3	Singin' in the Rain	US	1952	6	Gene Kelly, Debbie Reynolds, Donald O'Connor
4	Rear Window	US	1954	1	James Stewart, Grace Kelly

Director ID	Director_ Firstname	Director_L astname	Born	Died
1	Alfred	Hitchcock	1899	1980
2	Orson	Welles	1915	1985
3	….	…	…	…

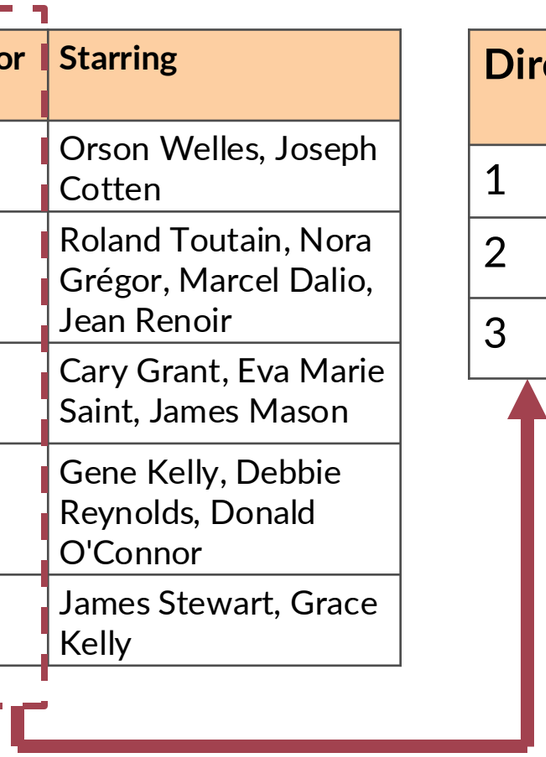

Link the director information between tables

Normalization

- Normal Form (NF)
  - 1NF: Simple attributes
  - 2NF: Attributes depend on the full key
  - 3NF: Non-key attributes do not depend on each other
  - And many others

	UNF (1970)	1NF (1970)	2NF (1971)	3NF (1971)	EKNF (1982)	BCNF (1974)	4NF (1977)	ETNF (2012)	5NF (1979)	DKNF (1981)	6NF (2003)
Primary key (no duplicate tuples)[4]	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Atomic columns (cells cannot have tables as values)[5]	✗	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Every non-trivial functional dependency either does not begin with a proper subset of a candidate key or ends with a prime attribute (no partial functional dependencies of non-prime attributes on candidate keys)[5]	✗	✗	✓	✓	✓	✓	✓	✓	✓	✓	✓
Every non-trivial functional dependency either begins with a superkey or ends with a prime attribute (no transitive functional dependencies of non-prime attributes on candidate keys)[5]	✗	✗	✗	✓	✓	✓	✓	✓	✓	✓	✓
Every non-trivial functional dependency either begins with a superkey or ends with an elementary prime attribute	✗	✗	✗	✗	✓	✓	✓	✓	✓	✓	N/A
Every non-trivial functional dependency begins with a superkey	✗	✗	✗	✗	✗	✓	✓	✓	✓	✓	N/A
Every non-trivial multivalued dependency begins with a superkey	✗	✗	✗	✗	✗	✗	✓	✓	✓	✓	N/A
Every join dependency has a superkey component[8]	✗	✗	✗	✗	✗	✗	✗	✓	✓	✓	N/A
Every join dependency has only superkey components	✗	✗	✗	✗	✗	✗	✗	✗	✓	✓	N/A
Every constraint is a consequence of domain constraints and key constraints	✗	✗	✗	✗	✗	✗	✗	✗	✗	✓	✗
Every join dependency is trivial	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✓

https://en.wikipedia.org/wiki/Database_normalization

Normalization

Every non key attribute must provide a fact about the key, the whole key, and nothing but the key.

William Kent (1936 – 2005)

William Kent. "A Simple Guide to Five Normal Forms in Relational Database Theory", Communications of the ACM 26 (2), Feb. 1983, pp. 120–125.

Normalization

Question:
What if there are multiple director?

- "First Norm Rule" (1NF)
  - Each column should only contain ONE piece of information

Movie ID	Movie Title	Country	Year	Director ID	Starring
0	Citizen Kane	US	1941	2	Orson Welles, Joseph Cotten
1	La règle du jeu	FR	1939	5	Roland Toutain, Nora Grégor, Marcel Dalio, Jean Renoir
2	North By Northwest	US	1959	1	Cary Grant, Eva Marie Saint, James Mason
3	Singin' in the Rain	US	1952	6	Gene Kelly, Debbie Reynolds, Donald O'Connor
4	Rear Window	US	1954	1	James Stewart, Grace Kelly

Director ID	Director_ Firstname	Director_Lastname	Born	Died
1	Alfred	Hitchcock	1899	1980
2	Orson	Welles	1915	1985
3	....	...	...	...

Link the director information between tables

Entity and Relationship

- A bad idea: Add more columns in the Movie table

Movie ID	Movie Title	Country	Year	Director ID	Director 2 ID	Director 3 ID	Starring

  - Waste of space (not too many movies have 3 directors, let alone 6)

Entity and Relationship

- A bad idea: Add more columns in the Movie table

Movie ID	Movie Title	Country	Year	Director ID	Director 2 ID	Director 3 ID	Starring

  - Waste of space (not too many movies have 3 directors, let alone 6)
  - How about starring? 10+ more columns?

Entity and Relationship

- Further refactoring of the tables …

Movie ID	Movie Title	Country	Year
0	Citizen Kane	US	1941
1	La règle du jeu	FR	1939
2	North By Northwest	US	1959
3	Singin' in the Rain	US	1952
4	Rear Window	US	1954

Movie Entities

Relationship?

Director ID	Director_ Firstname	Director_ Lastname	Born	Died
1	Alfred	Hitchcock	1899	1980
2	Orson	Welles	1915	1985
3	….	…	…	…

Director Entities

Entity and Relationship

- Further refactoring of the tables …

Movie ID	Movie Title	Country	Year
0	Citizen Kane	US	1941
1	La règle du jeu	FR	1939
2	North By Northwest	US	1959
3	Singin' in the Rain	US	1952
4	Rear Window	US	1954

Movie Entities

Directed By

Movie ID	Director ID
0	2
1	5
2	1

Relationship!

Director ID	Director_ Firstname	Director_ Lastname	Born	Died
1	Alfred	Hitchcock	1899	1980
2	Orson	Welles	1915	1985
3	….	…	…	…

Director Entities

Entity and Relationship

- Further refactoring of the tables …

Movie ID	Movie Title	Country	Year
0	Citizen Kane	US	1941
1	La règle du jeu	FR	1939
2	North By Northwest	US	1959
3	Singin' in the Rain	US	1952
4	Rear Window	US	1954

Movie Entities

Directed By

Movie ID	Director ID
0	2
1	5
2	1
…	…
16	8
16	9
16	10

Director ID	Director_ Firstname	Director_ Lastname	Born	Died
1	Alfred	Hitchcock	1899	1980
2	Orson	Welles	1915	1985
3	….	…	…	…

Director Entities

Question:
What if there are multiple director?

Answer:
Multiple rows in the relationship!

Entity and Relationship

- Starring -> Actor table
- Country -> Country and Region table
  - You can also link the movies with corresponding actors, countries/regions, etc.

- Entity Relationship Diagram (E/R Diagram, ER Diagram, ERD)
  - A way of representing entity tables and their relationships (relationship tables)

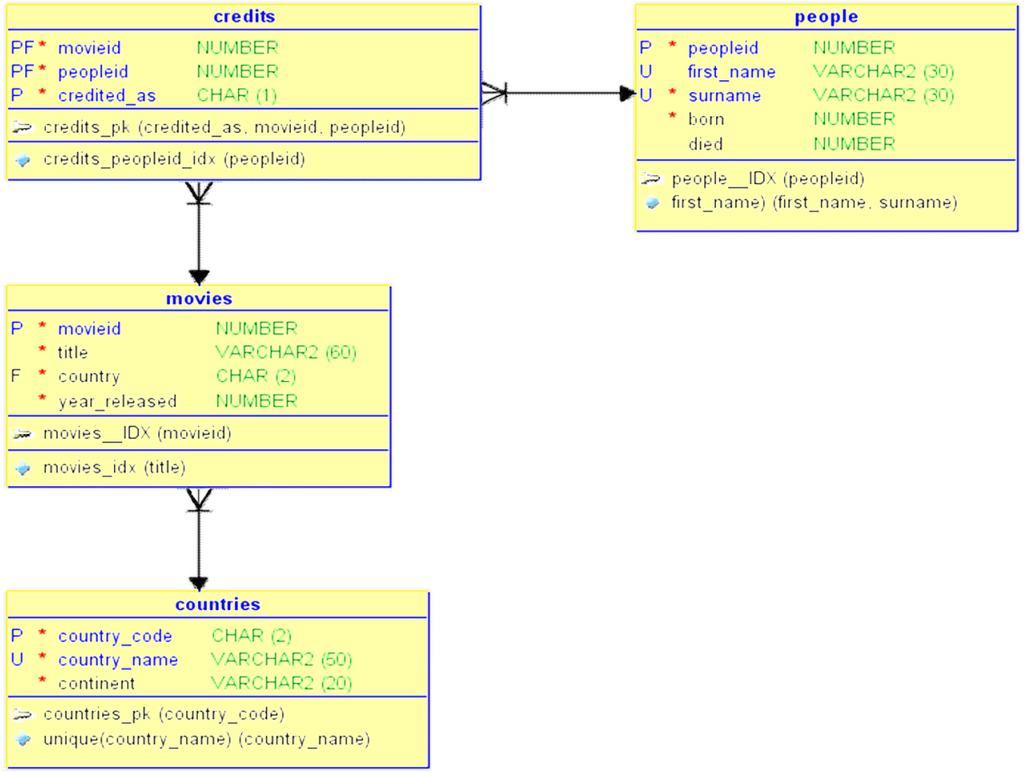

Outline

Part 1: Relational Language

Basic SQL
- Create Tables and Insert Data
- Basic Queries

Intermediate SQL
- Advanced Join
- Subquery
- Window Function

Advanced SQL
- Function
- Procedure
- Trigger
- View
- Transactions

Lab Only

Part 2: DB Design and Application

Relational DB Design
- The Entity-Relationship Model
- E-R Diagram Design

Relational Algebra

Normalization
- Functional Dependency
- Normal Forms
- FD Theory

Application Design
- Use Databases in Programs
- Complex Data Types (JSON, XML, etc.)

Part 3: Internal Mechanism Of Relational DB Systems

Storage Management
- Physical and Logical Storage Structure
- Indexing

Query Management
- Query Processing
- Query Optimization

Transaction Management
- Transaction
- Concurrency

Part 4: Miscellaneous Topics

Big Data
- Big Data Storage
- Query and Process Big Data

Advanced Topics
- Beyond Relational DB
- Data Analytics

Outline

- What we may not (fully) cover
  - Programming language support
    - We will focus on one or two languages
  - Parallel and distributed databases
  - Object-based databases
  - Blockchain
  - Advanced query processing and optimization techniques
  - Advanced relational algebra and calculus
  - Advanced data mining and analytics
  - Implementation details in relational DB

Data Definition and Manipulation

- Data Definition Language (DDL)
  - DDL compiler generates a set of table templates stored in a data dictionary
    - Database schema
    - Integrity constraints (primary key, etc.)
    - Authorization (who can access it)

```
create table lab(
    id serial primary key,
    address varchar(20) not null,
    time varchar(20) not null,
    capacity int,
    teacher varchar(20),
    unique (address,time)
);
```

Data Definition and Manipulation

- Data Manipulation Language (DML)
  - **Language** for accessing and updating the data organized by the appropriate data model (also known as query language)
- SQL (Structured Query Language)
  - Takes several tables as input (possibly only one) and always returns a single table

```
select * from lab;

select * from lab where time = '3-34';
```

Data Definition and Manipulation

- Data Manipulation Language (DML)
  - **Language** for accessing and updating the data organized by the appropriate data model (also known as query language)
- SQL (Structured Query Language)
  - Takes several tables as input (possibly only one) and always returns a single table

```
select * from lab;

select * from lab where time = '3-34';
```

Part 1: Warm-up Question (3 points)

(3 points, 1 points for each correct word) The lecturer has mentioned this meme multiple times in the classroom sessions, so let's try it again: What is "SQL" short for?

Database Users

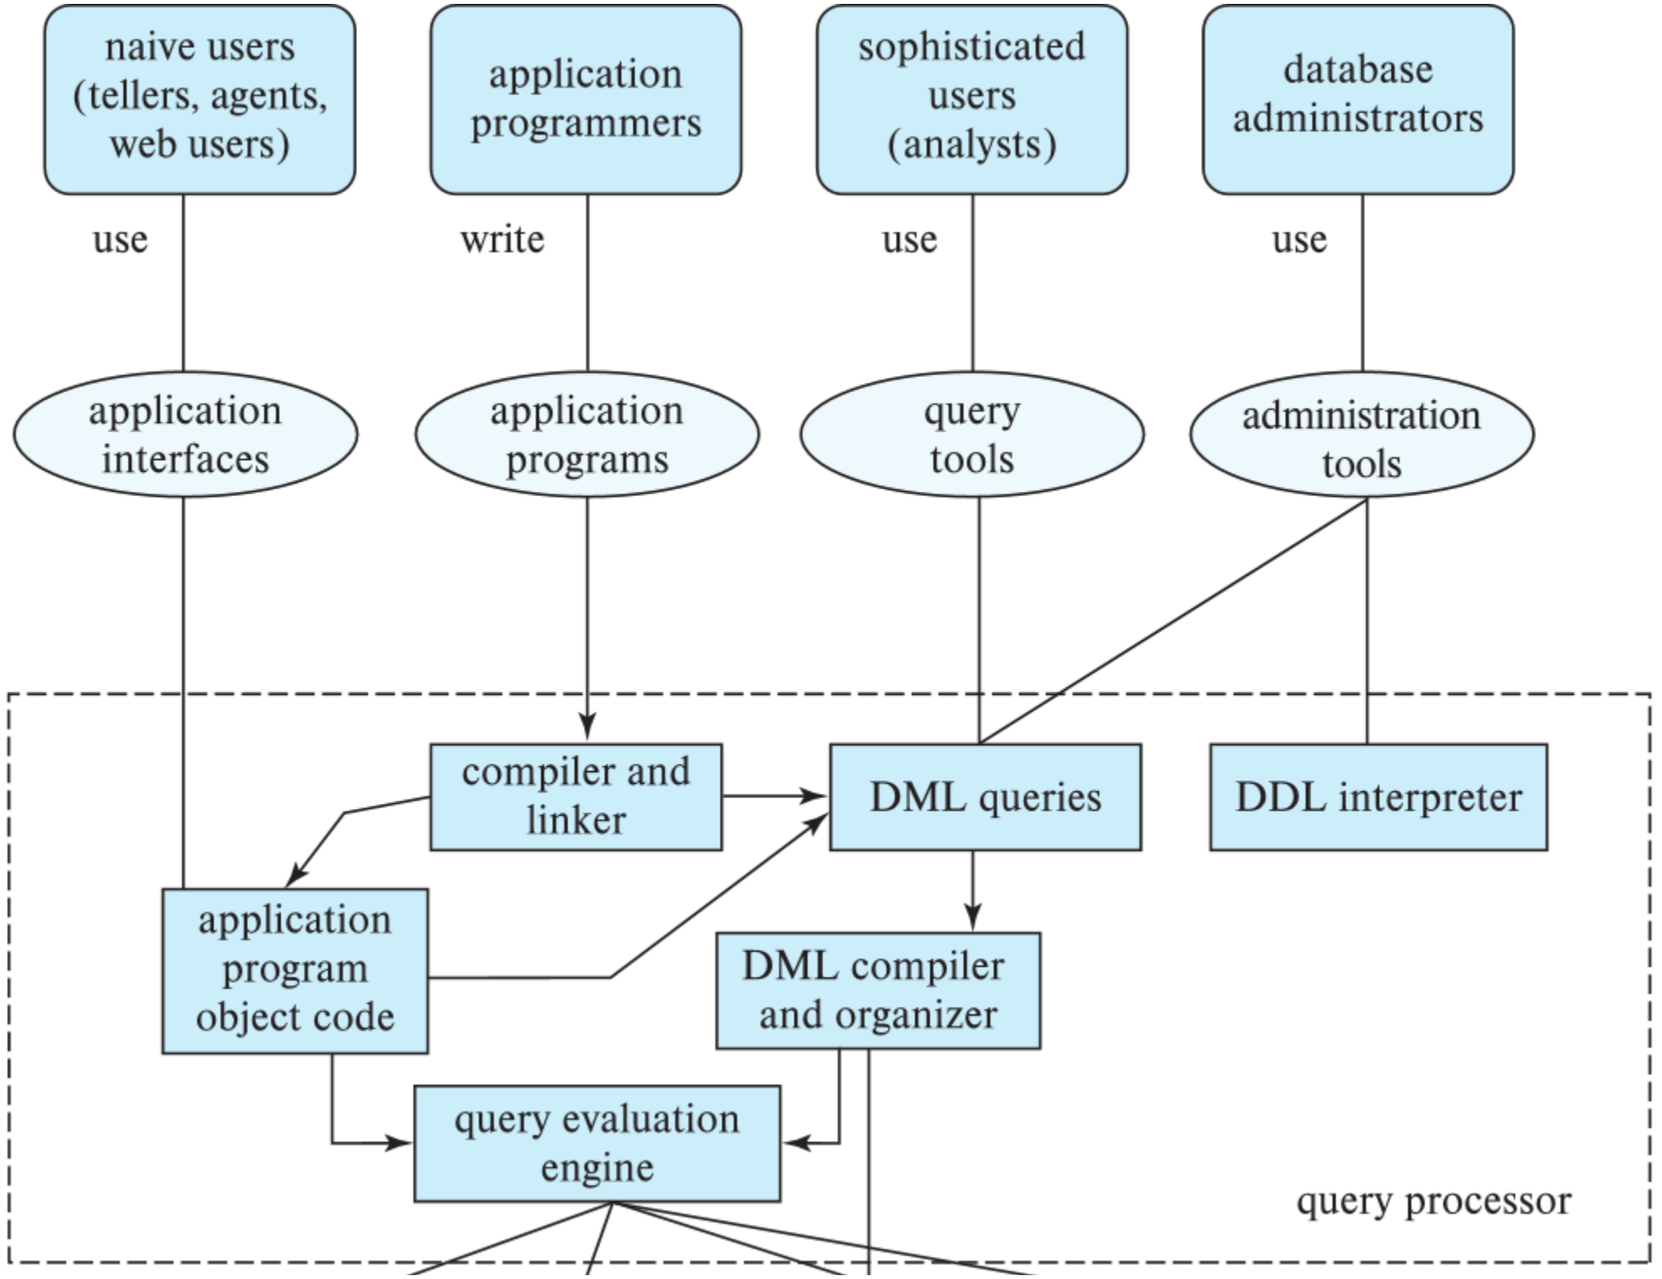

Database Architecture

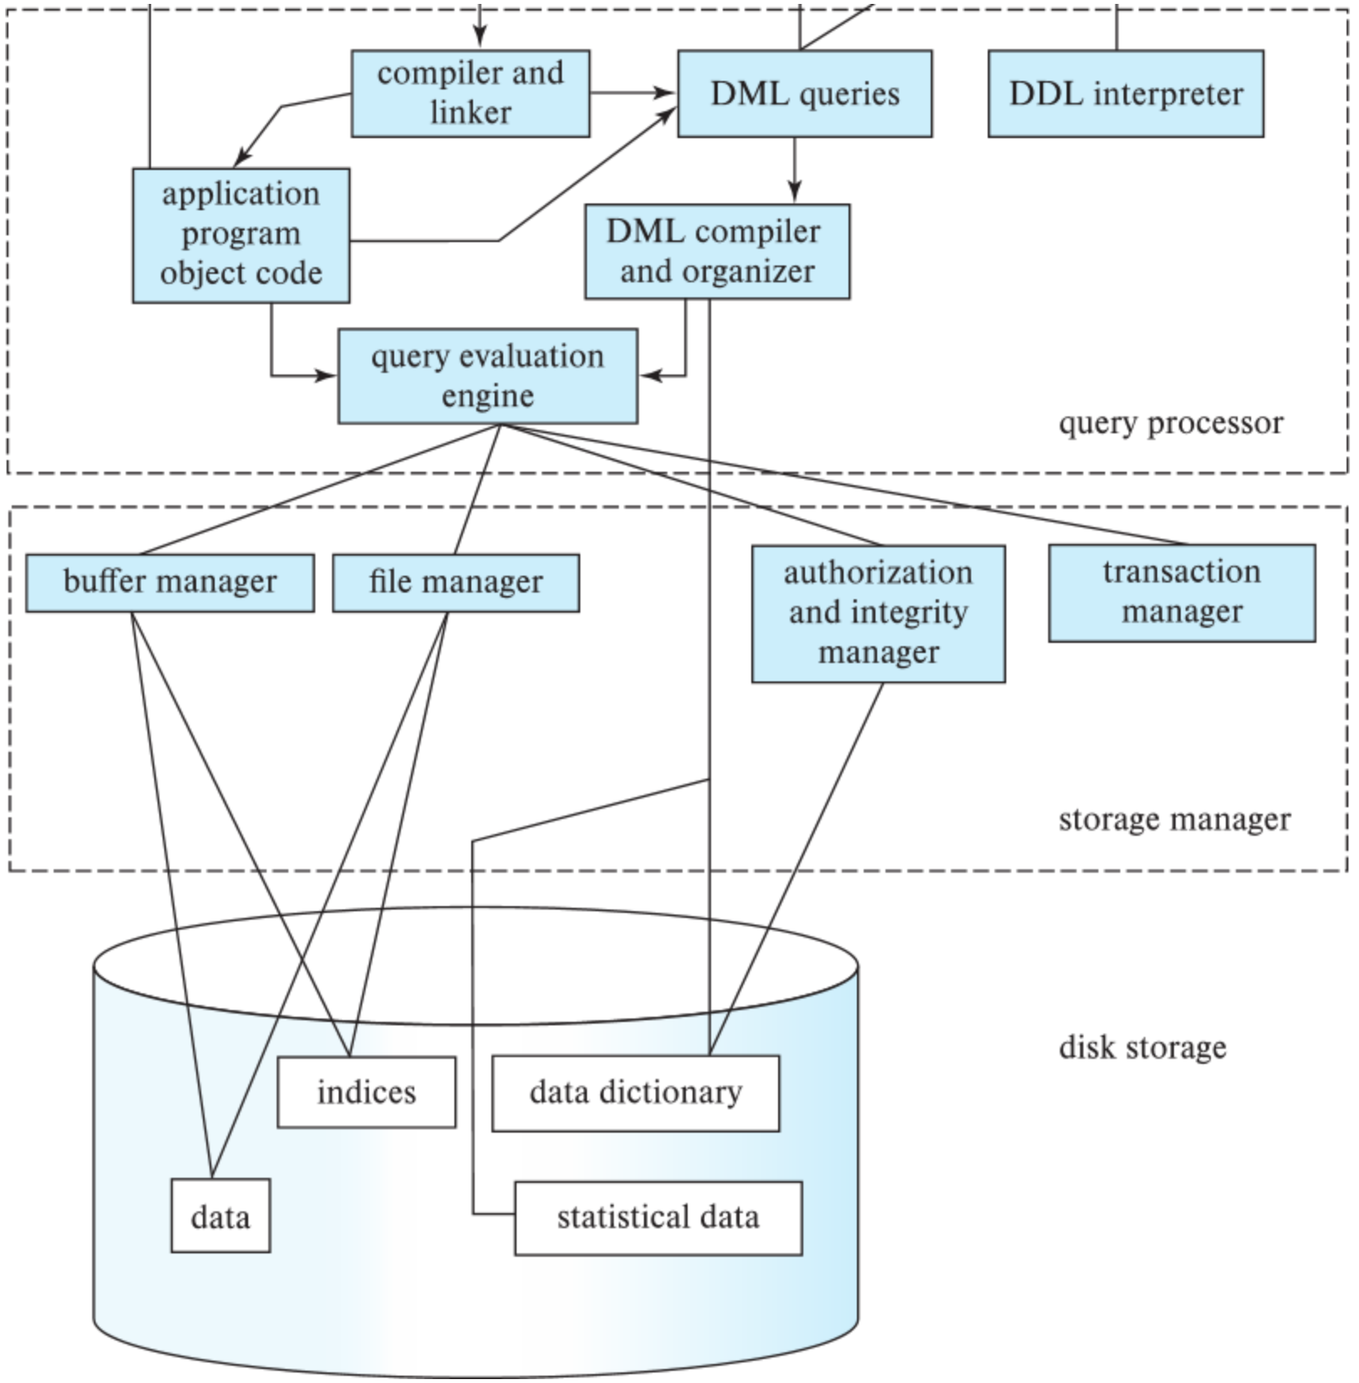

Database Applications

- Database applications are usually partitioned into two or three parts
- Application programs generally access databases through one of
  - Language extensions to allow embedded SQL
  - Application Program Interface (e.g., ODBC/JDBC) which allow SQL queries to be sent to a database system

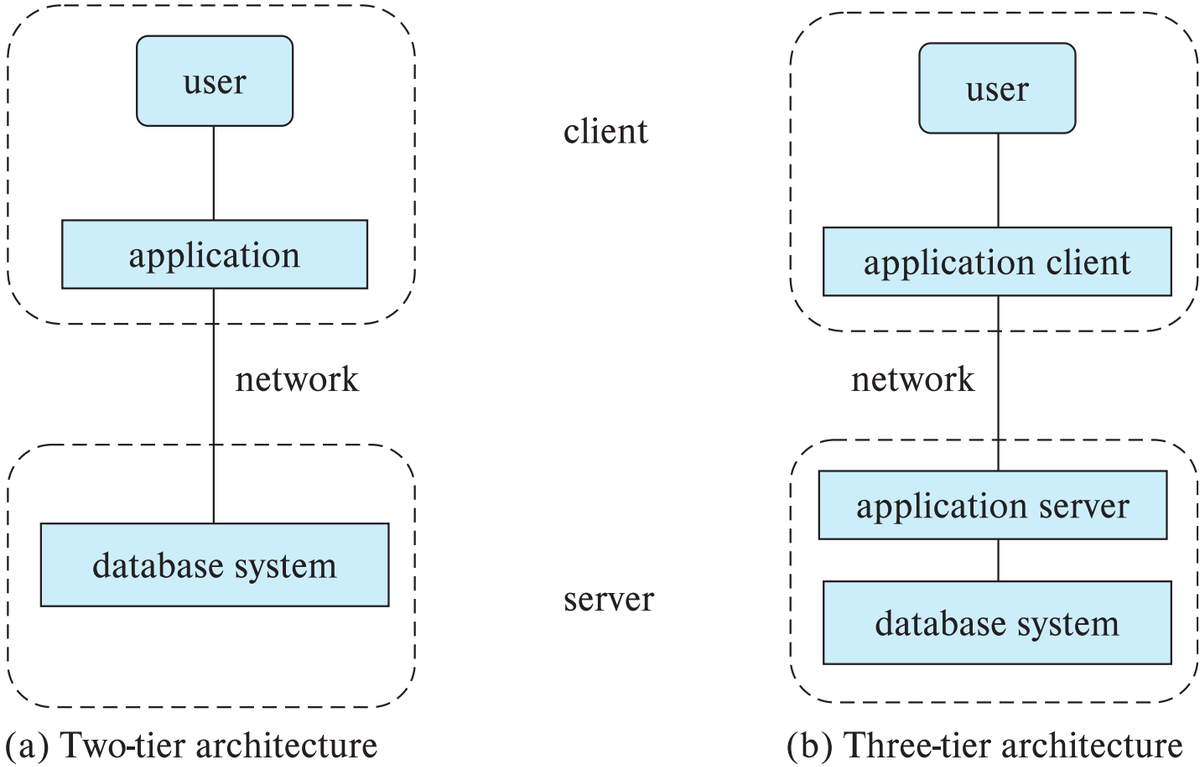

(a) Two-tier architecture

(b) Three-tier architecture

Some Important Prerequisite

Grading Policy

- Late Submission
  - We do not accept late submissions. All assigments, quizzes, and projects, etc. will receive a score of zero if you miss the deadline.
- Groups for Projects
  - Groups across different lab sessions are not allowed
  - Please try to find your teammate in the same lab session
- Grades
  - The teachers and TAs guarantee that your assignments and projects will be evaluated carefully and unbiasedly
  - We do not accept argueing with teachers over a certain grade once the decision has been made

Plagiarism

- Please read the regulations on academic misconduct on the Blackboard site
  - Blackboard – “Declaration Form Submission”
  - All documents are in the attachment of the assignment

- Upload the signed *Undergraduate Students Declaration Form* to Blackboard with the following file name format:
  - `SID_name.pdf`

  - * The submission entry will be announced later on Blackboard

Some Other Stuff

- Computing technologies advance very fast
  - Search online to learn more by yourself
    - Search engines (Google, Bing, Baidu, etc.), StackOverflow, GitHub.
  - The lecture notes can guide your self study

- You are encouraged to ask questions
  - At any time

- Practice makes perfect
  - No need to be afraid of trying new techniques/ideas/codes

Notices

- Find your own lab session (3 different groups)
  - Attendance is required
- WeCom Group
  - Teaching assistants will be there to help you
- (For international students) If you don’t know how to join the WeCom group or Blackboard site, please email me or find TAs as soon as possible
  - Contact Email: zhuym@sustech.edu.cn, mayx@sustech.edu.cn

Please remember to attend your lab class.
It is strongly recommended to install the required software before your lab session.