

# Principles of Database Systems (CS307)

## Lecture 13: Query Optimization

**Yuxin Ma**

Department of Computer Science and Engineering  
Southern University of Science and Technology

- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7<sup>th</sup> Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.

# Query Optimization

- Purpose of query optimization
  - Select an effective way to retrieve the data based on queries while spending the least computational effort
    - However, it is only “spending less computational effort” in most scenarios, not least

# Query Optimization

- Users don't need to consider the best way of writing queries
- Automated optimization can perform better (for most of the time)
  - Utilize the data dictionary
  - Real-time utilization based on physical storage changes
  - Optimizer can evaluate hundreds of execution plans in a very short time compared with human programmers
  - Human users do not need to learn advanced optimization techniques any more, which is conducted by optimizers instead

# An Example in the Movie Dataset

- The same query can be represented in different plans
  - E.g., retrieve the titles of those movies from China



```
select m.title  
from movies m, countries c  
where m.country = c.country_code and c.country_name = 'China';
```

# An Example in the Movie Dataset

- The corresponding relational algebra expressions:

$$(1) \Pi_{title} (\sigma_{movies.country = countries.country\_code \wedge countries.country = "China"}(movies \times countries))$$

$$(2) \Pi_{title} (\sigma_{countries.country = "China"}(movies \bowtie_{movies.country = countries.country\_code} countries))$$

$$(3) \Pi_{title} (movies \bowtie_{movies.country = countries.country\_code} (\sigma_{countries.country = "China"}(countries)))$$

# An Example in the Movie Dataset

- The corresponding relational algebra expressions:

(1)  $\Pi_{title} (\sigma_{movies.country = countries.country\_code \wedge countries.country = "China"}(movies \times countries))$

(2)  $\Pi_{title} (\sigma_{countries.country = "China"}(movies \bowtie_{movies.country = countries.country\_code} countries))$

(3)  $\Pi_{title} (movies \bowtie_{movies.country = countries.country\_code} (\sigma_{countries.country = "China"}(countries)))$

- In (1), a **full Cartesian product** will be computed, which costs huge time for matching all pairs and massive temporary storage space for the intermediate product table
- In (2), a **smaller intermediate join table** is to be cached, which saves some space
- In (3), **the filter**  $(\sigma_{c.country = "China"})$  **reduces the size of the right table** in the join operation, which saves a lot of time for pair matching and caching intermediate join table

# An Example in the Movie Dataset

- In addition, the filter operation can be further accelerated once an index is built upon the *country* column


(3)  $\Pi_{title} (\text{movies} \bowtie_{\text{movies.country} = \text{countries.country\_code}} (\sigma_{\text{countries.country} = \text{"China"}} (\text{countries})))$

# Generating Equivalent Expressions


- Alternative ways of evaluating a given query
  - Equivalent expressions
  - Different algorithms for each operation



```
select name, title
from instructor
  natural join (teaches natural join course)
where dept_name = 'Music' and year = 2009;
```



(a) Initial expression tree



(b) Tree after multiple transformations

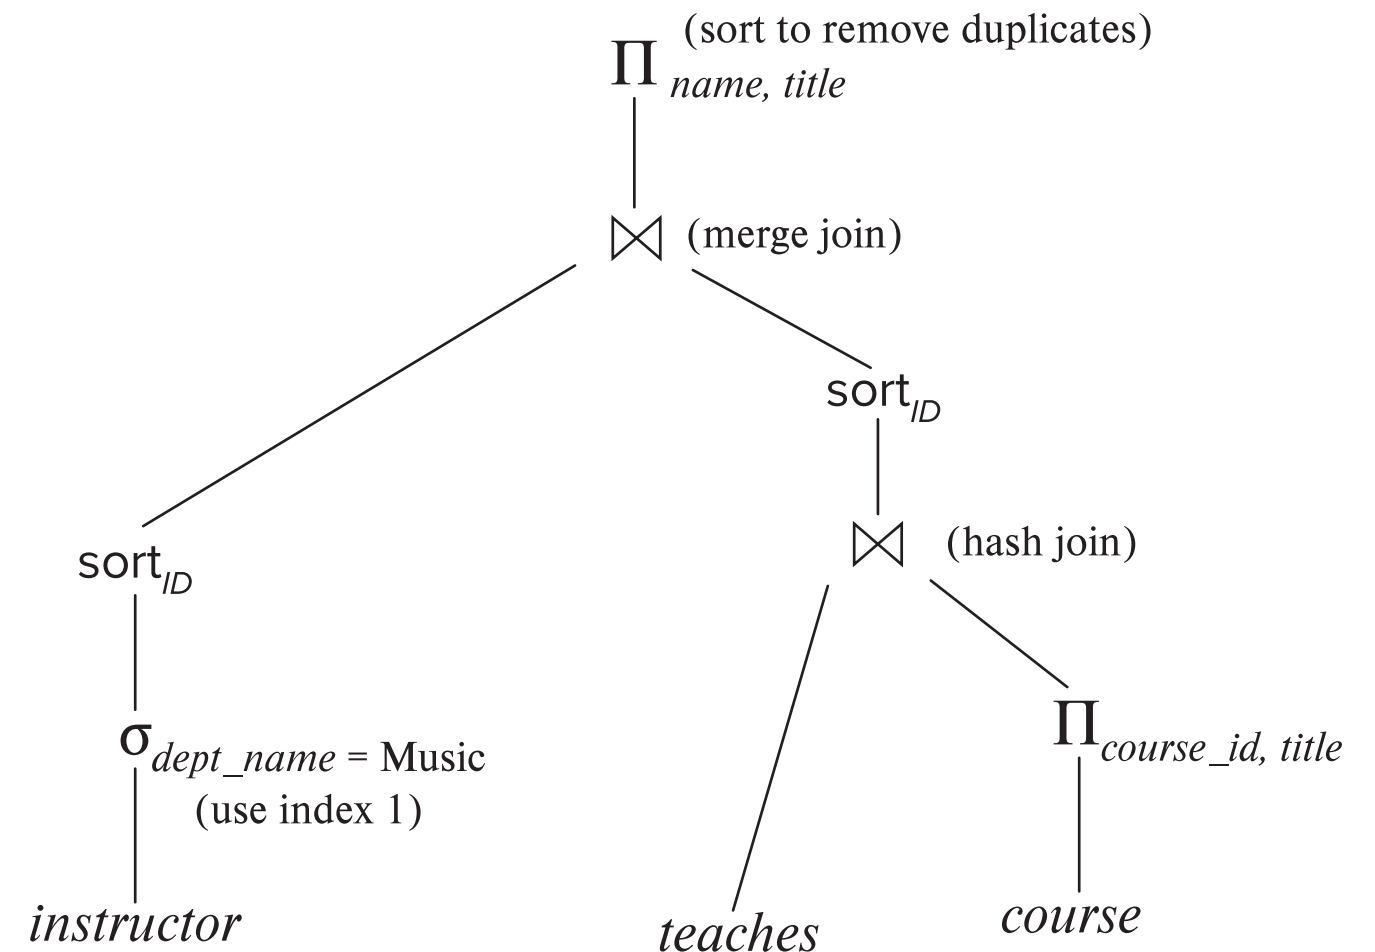


# Generating Equivalent Expressions

- An **evaluation plan** defines exactly what algorithm is used for each operation, and how the execution of the operations is coordinated



```
select name, title
from instructor
      natural join (teaches natural join course)
where dept_name = 'Music' and year = 2009;
```



# Transformation of Relational Expressions

- Two relational algebra expressions are said to be **equivalent** if the two expressions generate the same set of tuples on every legal database instance
  - Note: order of tuples is irrelevant
  - We don't care if they generate different results on databases that violate integrity constraints
- An **equivalence rule** says that expressions of two forms are equivalent
  - ... i.e., we can replace expression of the first form by second, or vice versa

# Equivalence Rules

1. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) \equiv \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. Selection operations are commutative

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) \equiv \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. Only the last in a sequence of projection operations is needed, the others can be omitted

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) \equiv \Pi_{L_1}(E)$$

where  $L_1 \subseteq L_2 \dots \subseteq L_n$

4. Selections can be combined with Cartesian products and theta joins

a)  $\sigma_{\theta}(E_1 \times E_2) \equiv E_1 \bowtie_{\theta} E_2$

b)  $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) \equiv E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

# Equivalence Rules

5. Theta-join operations (and natural joins) are commutative.

$$E_1 \bowtie E_2 \equiv E_2 \bowtie E_1$$

6. (a) Natural join operations are associative:

$$(E_1 \bowtie E_2) \bowtie E_3 \equiv E_1 \bowtie (E_2 \bowtie E_3)$$


(b) Theta joins are associative in the following manner:

$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 \equiv E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$



where  $\theta_2$  involves attributes from only  $E_2$  and  $E_3$

# Equivalence Rules



- Representation of Rule 5, 6(a) and 6(b) with diagrams



Rule 5




Rule 6.a



Rule 7.a

If  $\theta$  only has attributes from  $E_1$



# Equivalence Rules

7. The selection operation distributes over the theta join operation under the following two conditions:

- (a) When all the attributes in  $\theta_0$  involve only the attributes of one of the expressions ( $E_1$ ) being joined:

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) \equiv (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

- (b) When  $\theta_1$  involves only the attributes of  $E_1$  and  $\theta_2$  involves only the attributes of  $E_2$ :

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) \equiv (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$

# Equivalence Rules

8. The projection operation distributes over the theta join operation as follows:

(a) If  $\theta$  involves only attributes from  $L_1 \cup L_2$ :

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) \equiv \Pi_{L_1}(E_1) \bowtie_{\theta} \Pi_{L_2}(E_2)$$

(b) In general, consider a join  $E_1 \bowtie_{\theta} E_2$ :

- Let  $L_1$  and  $L_2$  be sets of attributes from  $E_1$  and  $E_2$ , respectively,
- Let  $L_3$  be attributes of  $E_1$  that are involved in join condition  $\theta$ , but are not in  $L_1 \cup L_2$ , and,
- Let  $L_4$  be attributes of  $E_2$  that are involved in join condition  $\theta$ , but are not in  $L_1 \cup L_2$ :

$$\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) \equiv \Pi_{L_1 \cup L_2}(\Pi_{L_1 \cup L_3}(E_1) \bowtie_{\theta} \Pi_{L_2 \cup L_4}(E_2))$$

\* Similar equivalences hold for left, right, and full outer join operations:  $\bowtie$ ,  $\ltimes$ , and  $\Join$

# Equivalence Rules

9. The set operations union and intersection are commutative

$$E_1 \cup E_2 \equiv E_2 \cup E_1$$

$$E_1 \cap E_2 \equiv E_2 \cap E_1$$

- However, set difference is not commutative

10. Set union and intersection are associative

$$(E_1 \cup E_2) \cup E_3 \equiv E_1 \cup (E_2 \cup E_3)$$

$$(E_1 \cap E_2) \cap E_3 \equiv E_1 \cap (E_2 \cap E_3)$$



# Equivalence Rules

11. The selection operation distributes over  $\cup$ ,  $\cap$  and  $-$

(a)  $\sigma_{\theta} (E_1 \cup E_2) \equiv \sigma_{\theta} (E_1) \cup \sigma_{\theta}(E_2)$

(b)  $\sigma_{\theta} (E_1 \cap E_2) \equiv \sigma_{\theta} (E_1) \cap \sigma_{\theta}(E_2)$

(c)  $\sigma_{\theta} (E_1 - E_2) \equiv \sigma_{\theta} (E_1) - \sigma_{\theta}(E_2)$

(d)  $\sigma_{\theta} (E_1 \cap E_2) \equiv \sigma_{\theta}(E_1) \cap E_2$

(e)  $\sigma_{\theta} (E_1 - E_2) \equiv \sigma_{\theta}(E_1) - E_2$

- \* The preceding equivalence does not hold for  $\cup$

12. The projection operation distributes over union

$$\Pi_L(E_1 \cup E_2) \equiv (\Pi_L(E_1)) \cup (\Pi_L(E_2))$$

# Transformation Example: Pushing Selections

- Query: Find the names of all instructors in the Music department, along with the titles of the courses (in the Music department) that they teach

```
select name, title
from instructor natural join (teaches natural join course
where dept_name = 'Music');
```

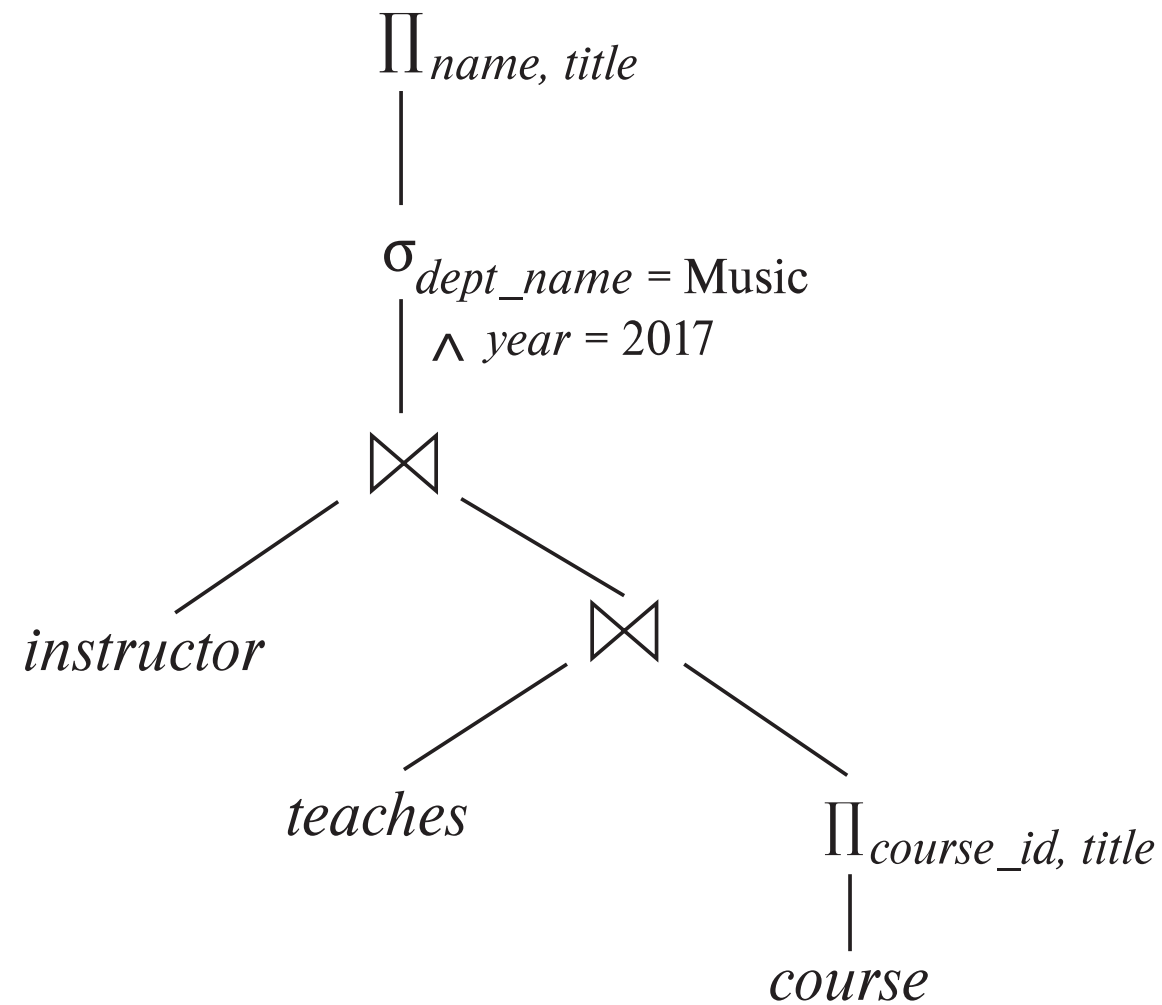
- $\Pi_{name, title}(\sigma_{dept\_name = 'Music'}(instructor \bowtie (teaches \bowtie \Pi_{course\_id, title}(course))))$
- Transformation using rule 7(a):
  - $\Pi_{name, title}((\sigma_{dept\_name = 'Music'}(instructor)) \bowtie (teaches \bowtie \Pi_{course\_id, title}(course)))$

Performing the selection as early as possible reduces the size of the relation to be joined

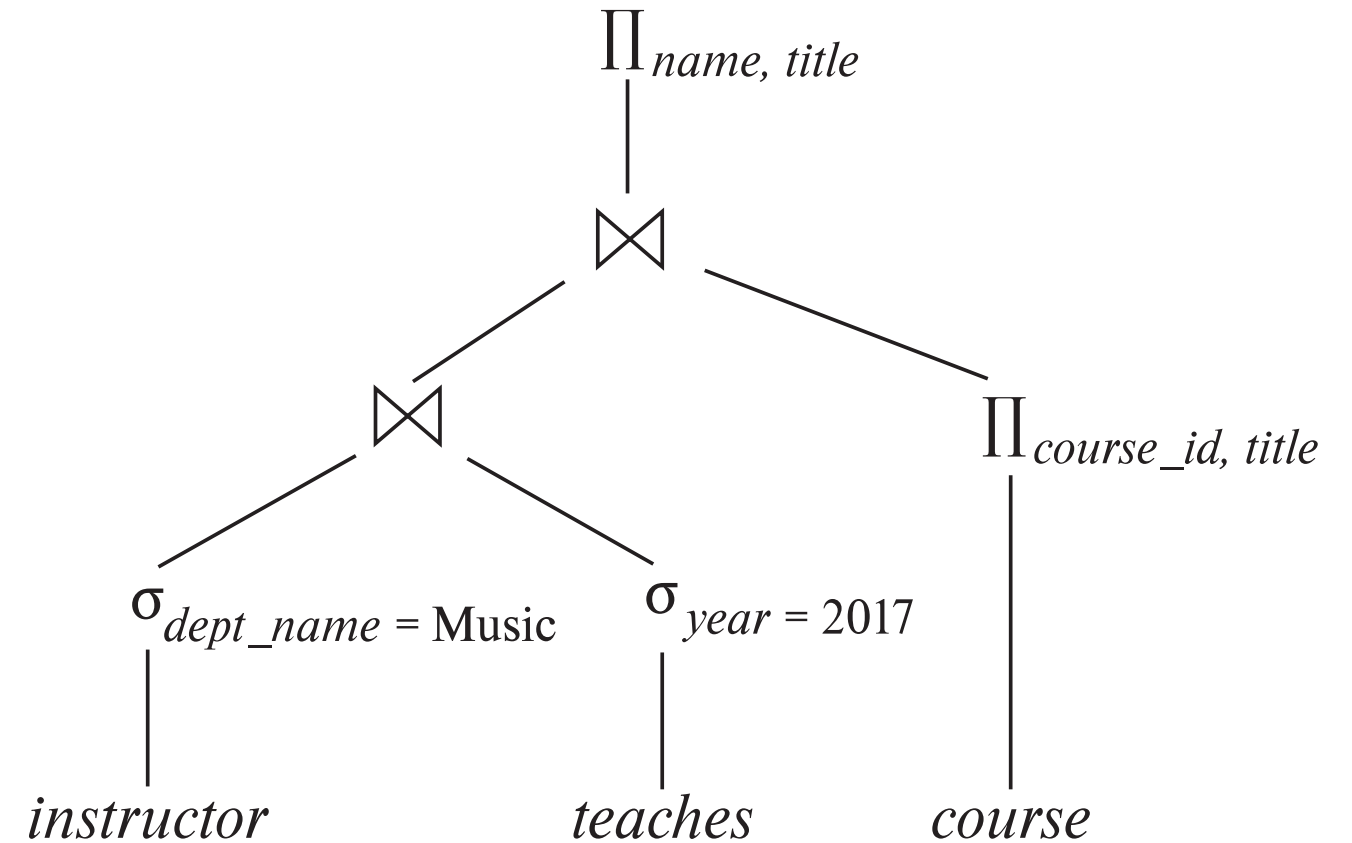
# Transformation Example: Multiple Transformations

- Query: Find the names of all instructors in the Music department who have taught a course in 2017, along with the titles of the courses that they taught
  - $\Pi_{name, title}(\sigma_{dept\_name = \text{"Music"} \wedge year = 2017} (instructor \bowtie (teaches \bowtie \Pi_{course\_id, title} (course))))$
- Transformation using join associatively (Rule 6(a)):
  - $\Pi_{name, title}(\sigma_{dept\_name = \text{"Music"} \wedge year = 2017} ((instructor \bowtie teaches) \bowtie \Pi_{course\_id, title} (course)))$
- Second form provides an opportunity to apply the “perform selections early” rule, resulting in the subexpression:
  - $\sigma_{dept\_name = \text{"Music"}} (instructor) \bowtie \sigma_{year = 2017} (teaches)$

# Transformation Example: Multiple Transformations



(a) Initial expression tree



(b) Tree after multiple transformations

## \* Transformation Example: Pushing Projections

- Consider  $\Pi_{name, title} (\sigma_{dept\_name = \text{"Music"}} (instructor) \bowtie teaches) \bowtie \Pi_{course\_id, title} (course))$

- When we compute

$$(\sigma_{dept\_name = \text{"Music"}} (instructor) \bowtie teaches),$$

we obtain a relation whose schema is:

$(ID, name, dept\_name, salary, course\_id, sec\_id, semester, year)$

- Push projections using equivalence rules 8a and 8b; eliminate unneeded attributes from intermediate results to get:

$$\Pi_{name, title} (\Pi_{name, course\_id} (\sigma_{dept\_name = \text{"Music"}} (instructor) \bowtie teaches) \bowtie \Pi_{course\_id, title} (course)))$$

Performing the projections as early as possible reduces the size of the relation to be joined

# Join Ordering Example

- For all relations  $r_1, r_2$ , and  $r_3$ ,

$$(r_1 \bowtie r_2) \bowtie r_3 = r_1 \bowtie (r_2 \bowtie r_3)$$

- \* (Join Associativity)  $\bowtie$
- If  $r_2 \bowtie r_3$  is quite large and  $r_1 \bowtie r_2$  is small, we choose
$$(r_1 \bowtie r_2) \bowtie r_3$$
so that we compute and store a smaller temporary relation

# Cost Estimation

- Cost difference between evaluation plans for a query can be enormous
  - E.g., seconds vs. days in some cases
- Steps in cost-based query optimization
  - 1. Generate logically equivalent expressions using equivalence rules
  - 2. Annotate resultant expressions to get alternative query plans
  - 3. Choose the cheapest plan based on estimated cost

# Cost Estimation

- Estimation of plan cost based on:
  - Statistical information about relations, such as:
    - number of tuples, number of distinct values for an attribute
  - Statistics estimation for intermediate results
    - to compute cost of complex expressions
  - Cost formulae for algorithms, computed using statistics

For more, please refer to Section 16.3 “Estimating Statistics of Expression Results” in the reference textbook