



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

# CS323 Lab 2

Yepang Liu

[liuyp1@sustech.edu.cn](mailto:liuyp1@sustech.edu.cn)

# Agenda

- Recognizing tokens using transition diagram
- Introduction to Flex

# Recognition of Tokens

- Lexical analyzer examines the input string and finds a prefix that matches one of the token patterns
- The first thing when building a lexical analyzer is to define the patterns of tokens using regular definitions
- A special token:  $\text{ws} \rightarrow (\text{blank} \mid \text{tab} \mid \text{newline})^+$ 
  - When the lexical analyzer recognizes a whitespace token, it does not return it to the parser, but restarts from the next character

# Example: Patterns and Tokens

<i>digit</i>	$\rightarrow$	[0-9]
<i>digits</i>	$\rightarrow$	<i>digit</i> <sup>+</sup>
<i>number</i>	$\rightarrow$	<i>digits</i> ( . <i>digits</i> )? ( E [+-]? <i>digits</i> )?
<i>letter</i>	$\rightarrow$	[A-Za-z]
<i>id</i>	$\rightarrow$	<i>letter</i> ( <i>letter</i>   <i>digit</i> )*
<i>if</i>	$\rightarrow$	if
<i>then</i>	$\rightarrow$	then
<i>else</i>	$\rightarrow$	else
<i>relop</i>	$\rightarrow$	<   >   <=   >=   =   <>

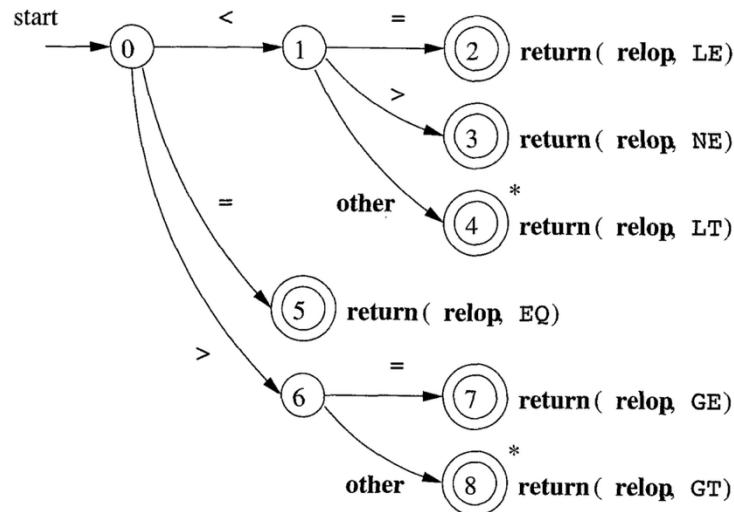
Patterns for tokens

LEXEMES	TOKEN NAME	ATTRIBUTE VALUE
Any ws	-	-
if	if	-
then	then	-
else	else	-
Any id	id	Pointer to table entry
Any number	number	Pointer to table entry
<	relop	LT
<=	relop	LE
=	relop	EQ
<>	relop	NE
>	relop	GT
>=	relop	GE

Lexemes, tokens, and attribute values

# Transition Diagrams (状态转换图)

- An important step in constructing a lexical analyzer is to convert patterns into “**transition diagrams**”
- Transition diagrams have a collection of nodes, called *states* (状态) and *edges* (边) directed from one node to another

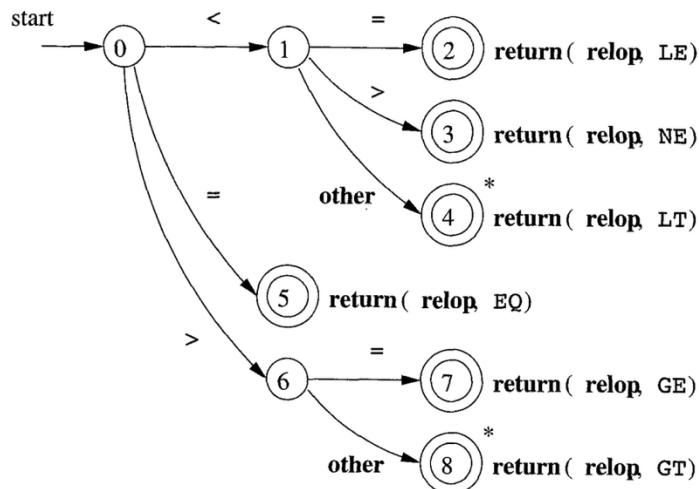


LEXEMES	TOKEN NAME	ATTRIBUTE VALUE
<	relop	LT
<=	relop	LE
=	relop	EQ
<>	relop	NE
>	relop	GT
>=	relop	GE

The transition diagram in the left recognizes **relop** tokens

# States

- Represent conditions that could occur during the process of scanning (i.e., what characters we have seen)
- The *start state* (开始状态), or *initial state*, is indicated by an edge labeled “start”, which enters from nowhere
- Certain states are said to be *accepting* (接受状态), or *final*, indicating that a lexeme has been found

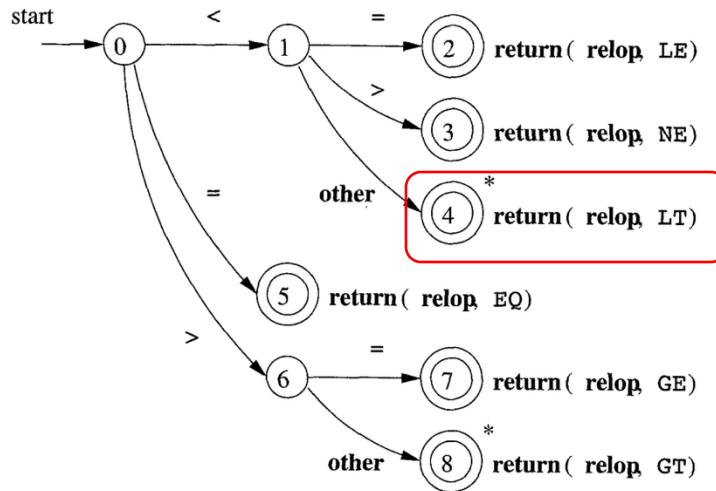


States 2-8 are accepting. They return a pair (token name, attribute value).

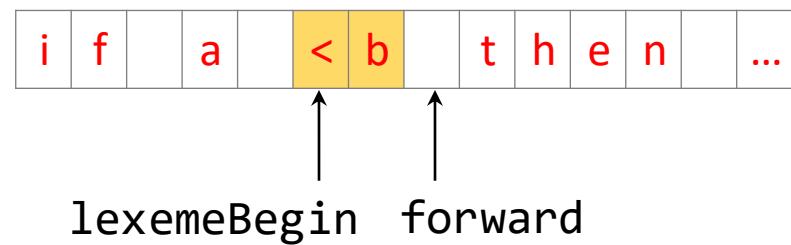
By convention, we indicate accepting states by **double circles**

# The Retract Action

- At certain accepting states, the found lexeme may not contain all characters that we have seen from the start state (such states are annotated with \*)
- When entering \* states, it is necessary to **retract** (回退) the **forward** pointer, which points to the next char in the input string



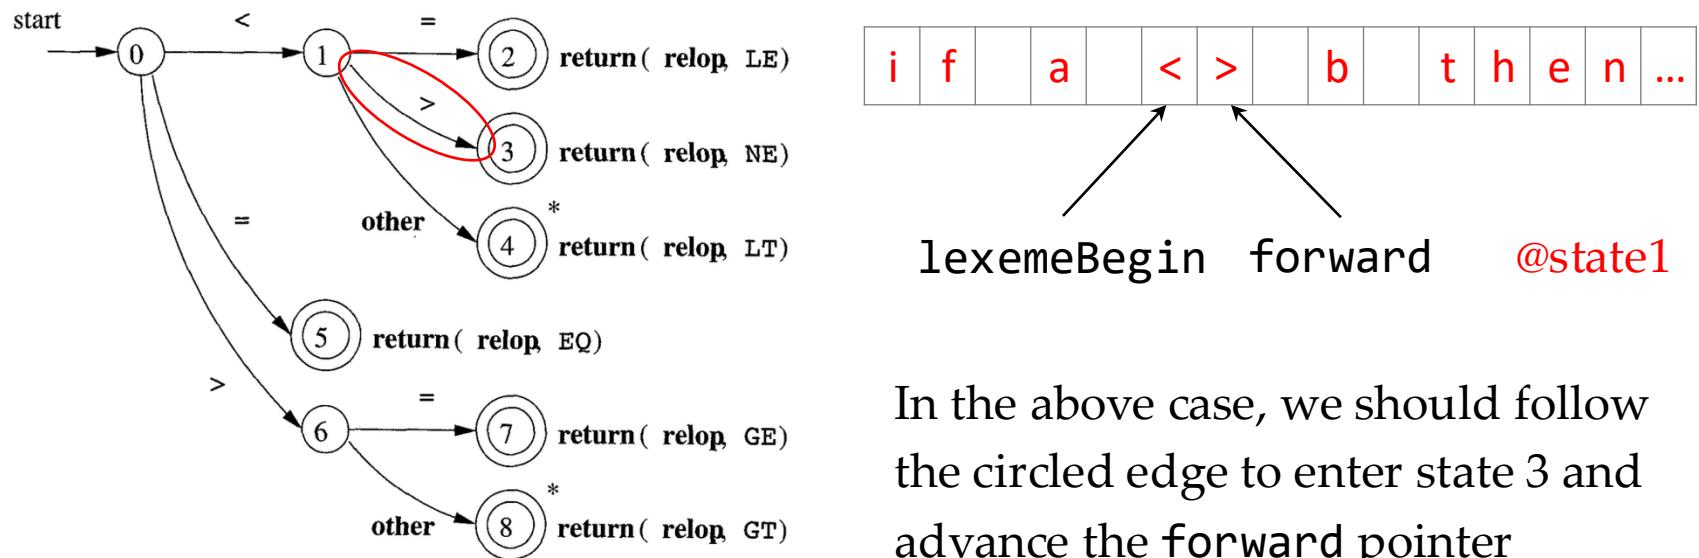
- The found lexeme: <
- The characters we've seen: <b



We should retract **forward** one step back

# Edges

- *Edges* are directed from one state to another
- Each edge is labeled by a symbol or set of symbols

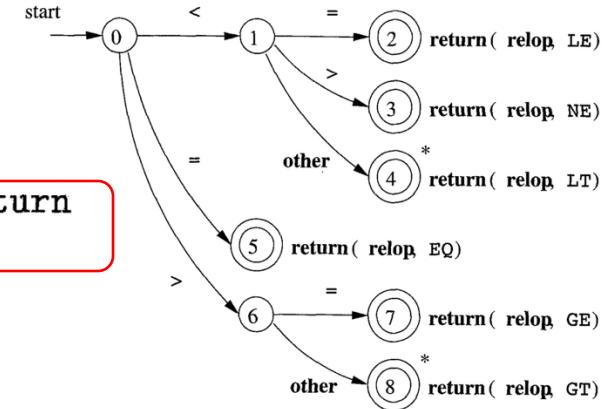


In the above case, we should follow the circled edge to enter state 3 and advance the **forward** pointer

# Building a Lexical Analyzer from Transition Diagrams

```
TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
        or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                if ( c == '<' ) state = 1;
                else if ( c == '=' ) state = 5;
                else if ( c == '>' ) state = 6;
                else fail(); /* lexeme is not a relop */
                break;
            case 1: ...
            ...
            case 8: retract();
                retToken.attribute = GT;
                return(retToken);
        }
    }
}
```

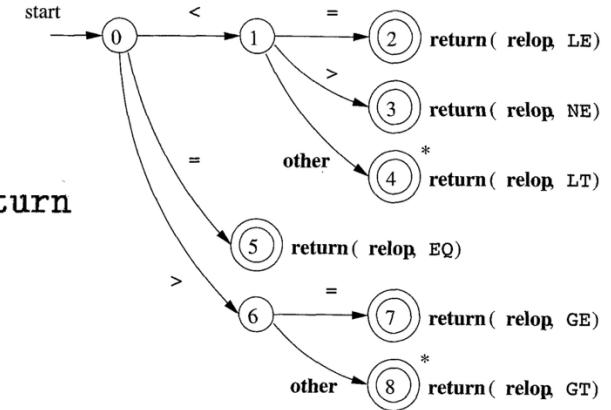
Sketch implementation of relop transition diagram



# Building a Lexical Analyzer from Transition Diagrams

```
TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
        or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                if ( c == '<' ) state = 1;
                else if ( c == '=' ) state = 5;
                else if ( c == '>' ) state = 6;
                else fail(); /* lexeme is not a relop */
                break;
            case 1: ...
            ...
            case 8: retract();
                retToken.attribute = GT;
                return(retToken);
        }
    }
}
```

Sketch implementation of relop transition diagram

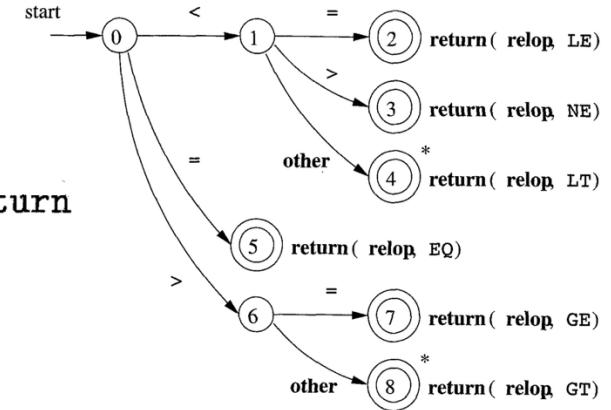


Use a variable **state** to record  
the current state

# Building a Lexical Analyzer from Transition Diagrams

```
TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
        or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                if ( c == '<' ) state = 1;
                else if ( c == '=' ) state = 5;
                else if ( c == '>' ) state = 6;
                else fail(); /* lexeme is not a relop */
                break;
            case 1: ...
            ...
            case 8: retract();
                retToken.attribute = GT;
                return(retToken);
        }
    }
}
```

Sketch implementation of relop transition diagram

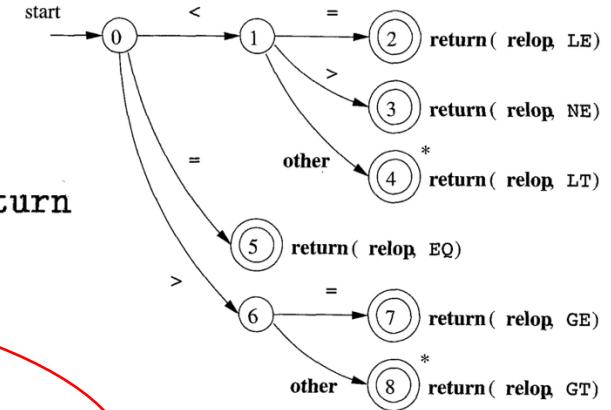


A **switch** statement based on the value of **state** takes us to the processing code

# Building a Lexical Analyzer from Transition Diagrams

```
TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
        or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                if ( c == '<' ) state = 1;
                else if ( c == '=' ) state = 5;
                else if ( c == '>' ) state = 6;
                else fail(); /* lexeme is not a relop */
                break;
            case 1: ...
            ...
            case 8: retract();
                retToken.attribute = GT;
                return(retToken);
        }
    }
}
```

Sketch implementation of relop transition diagram



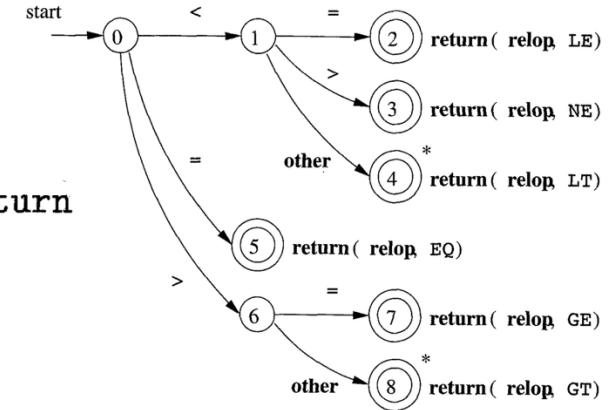
The code of a normal state:

1. Read the next character
2. Determine the next state
3. If step 2 fails, do error recovery

# Building a Lexical Analyzer from Transition Diagrams

```
TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
        or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                if ( c == '<' ) state = 1;
                else if ( c == '=' ) state = 5;
                else if ( c == '>' ) state = 6;
                else fail(); /* lexeme is not a relop */
                break;
            case 1: ...
            ...
            case 8: retract();
                retToken.attribute = GT;
                return(retToken);
        }
    }
}
```

Sketch implementation of relop transition diagram



The code of an accepting state:

1. Perform retraction if the state has \*
2. Set token attribute values
3. Return the token to parser

# Building the Entire Lexical Analyzer

- **Strategy 1:** Try the transition diagram for each type of token sequentially
  - `fail()` resets the pointer forward and tries the next diagram
- **Problem:** Not efficient
  - May need to try many irrelevant diagrams whose first edge does not match the first character in the input stream

# Building the Entire Lexical Analyzer

- **Strategy 2:** Run transition diagrams in parallel
  - Need to resolve the case where one diagram finds a lexeme and others are still able to process input.
  - **Solution:** take the longest prefix of the input that matches any pattern
- **Problem:** Requires special hardware for parallel simulation, may degenerate into the sequential strategy on certain machines

# Building the Entire Lexical Analyzer

- **Strategy 3:** Combining all transition diagrams into one
  - Allow the transition diagram to read input until there is no possible next state
  - Take the longest lexeme that matched any pattern
- This is **a commonly-adopted strategy** in real-world compiler implementation (efficient & requires no special hardware)



How? Be patient ☺, we will talk about this later.

# Agenda

- Recognizing tokens using transition diagram
- Introduction to Flex

# The Lexical-Analyzer Generator Lex

- Lex, or a more recent tool Flex, allows one to specify a lexical analyzer by specifying regexps to describe patterns for tokens
- Often used with Yacc/Bison to create the frontend of compiler

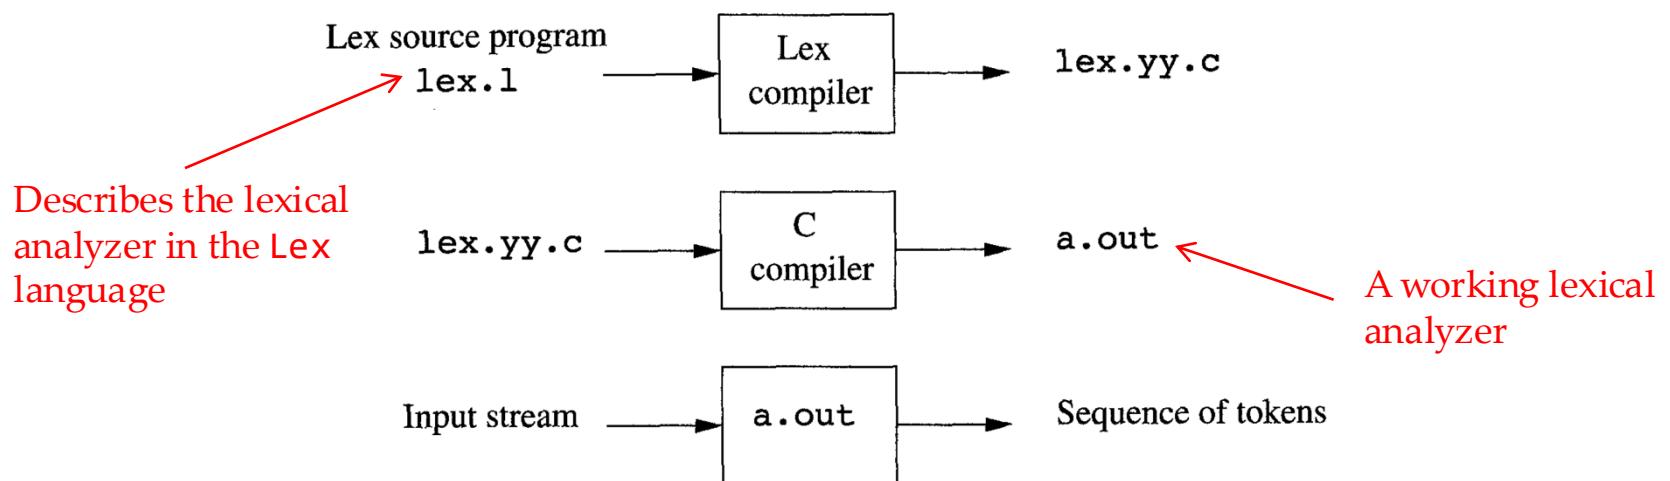


Figure 3.22: Creating a lexical analyzer with Lex

# Structure of Lex Programs

- A Lex program has three sections separated by %%
- Declaration (声明)
  - Variables, constants (e.g., token names)
  - Regular definitions
- Translation rules (转换规则) in the form “Pattern {Action}”
  - Each pattern (模式) is a regexp (may use the regular definitions of the declaration section)
  - Actions (动作) are fragments of code, typically in C, which are executed when the pattern is matched
- Auxiliary functions section (辅助函数)
  - Additional functions that can be used in the actions

# Lex Program Example

```
%{  
    /* definitions of manifest constants  
    LT, LE, EQ, NE, GT, GE,  
    IF, THEN, ELSE, ID, NUMBER, RELOP */  
}  
/* regular definitions */
```

```
delim      [ \t\n]  
ws         {delim}+  
letter     [A-Za-z]  
digit      [0-9]  
id          {letter}({letter}|{digit})*  
number     {digit}+(\.{digit}+)?(E[+-]?{digit}+)?
```

```
%%
```

Anything in between %{ and }% is copied directly to lex.yyy.c.  
In the example, there is only a comment, not real C code to define manifest constants

Regular definitions that can be used in translation rules

Section separator

# Lex Program Example Cont.

```
{ws}      /* no action and no return */  
if       {return(IF);}  
then     {return(THEN);}  
else     {return(ELSE);}  
{id}      yyval = (int) installID(); return(ID);  
{number}  yyval = (int) installNum(); return(NUMBER);  
"<"  
"<="  
"=="  
"<>"  
">"  
">=" %%  
A global variable that stores a pointer to the symbol table entry for the lexeme.  
Can be used by the parser or a later component of the compiler.
```

Literal strings\*

Continue to recognize other tokens

Return token name to the parser

Place the lexeme found in the symbol table

\* The characters inside have no special meaning (even if it is a special one such as \*).

# Lex Program Example Cont.

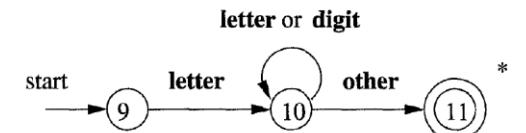
- Everything in the auxiliary function section is copied directly to the file `lex.yy.c`
- Auxiliary functions may be used in actions in the translation rules

```
int installID() /* function to install the lexeme, whose
                  first character is pointed to by yytext,
                  and whose length is yylen, into the
                  symbol table and return a pointer
                  thereto */
}
Variables defined and set automatically
by the lexical analyzer Lex generates

int installNum() /* similar to installID, but puts numer-
                  ical constants into a separate table */
}
```

# Conflict Resolution

- When the generated lexical analyzer runs, it analyzes the input looking for **prefixes that match any of its patterns.**\*
  - Rule 1:** If it finds multiple such prefixes, it takes the **longest** one
    - The analyzer will treat `<=` as a single lexeme, rather than `<` as one lexeme and `=` as the next
  - Rule 2:** If it finds a prefix matching different patterns, **the pattern listed first in the Lex program is chosen.**
    - Identifier pattern can also match keywords
    - If the keyword patterns are listed before identifier pattern, the lexical analyzer will not recognize keywords as identifiers



\* See Flex manual for details (Chapter 8: How the input is matched) at <http://dinosaur.compilertools.net/flex/>

# Flex

- Flex的前身是Lex。Lex是1975年由Mike Lesk和当时还在贝尔实验室做暑期实习的Eric Schmidt（前谷歌CEO），共同完成的一款基于Unix环境的词法分析程序生成工具。虽然Lex很出名并被广泛使用，但它的低效和诸多问题也使其颇受诟病。
- 1987年伯克利实验室（隶属美国能源部的国家实验室）的Vern Paxson使用C语言重写Lex，并将这个新程序命名为Flex（Fast Lexical Analyzer Generator）。无论从效率上还是稳定性上，Flex都远远好于它的前辈Lex。

\*我们在Linux下使用的是Flex在BSD License下的版本（和Bison不同，Flex不属于GNU计划）。

# An Example Flex Program

- A word-counting program (the Flex source file can be found on our blackboard course site)
- Build the program with the following commands
  - `flex lex.l`
    - After running the command, you will see a `lex.yy.c` file generated
  - `gcc lex.yy.c -lfl -o wc.out`
    - After running the command, you will see an executable file `wc.out` generated

Note: Install `gcc` first if it is not available on your machine: `sudo apt install gcc`

# An Example Flex Program

```
cs323@deb-cs323-compilers:~/cs323/lab2$ ls -l
total 12
-rw-r--r-- 1 cs323 cs323 6525 Sep 14 05:25 inferno3.txt
-rw-r--r-- 1 cs323 cs323 1242 Sep 14 05:25 lex.l
cs323@deb-cs323-compilers:~/cs323/lab2$ flex lex.l
cs323@deb-cs323-compilers:~/cs323/lab2$ ls -l
total 60
-rw-r--r-- 1 cs323 cs323 6525 Sep 14 05:25 inferno3.txt
-rw-r--r-- 1 cs323 cs323 1242 Sep 14 05:25 lex.l
-rw-rw-r-- 1 cs323 cs323 45322 Sep 14 05:42 lex.yy.c
cs323@deb-cs323-compilers:~/cs323/lab2$ gcc lex.yy.c -lfl -o wc.out
cs323@deb-cs323-compilers:~/cs323/lab2$ ls -l
total 88
-rw-r--r-- 1 cs323 cs323 6525 Sep 14 05:25 inferno3.txt
-rw-r--r-- 1 cs323 cs323 1242 Sep 14 05:25 lex.l
-rw-rw-r-- 1 cs323 cs323 45322 Sep 14 05:42 lex.yy.c
-rwxrwxr-x 1 cs323 cs323 27432 Sep 14 05:43 wc.out
cs323@deb-cs323-compilers:~/cs323/lab2$ ./wc.out inferno3.txt
#lines #words #chars file path
162    1088    6525  inferno3.txt
cs323@deb-cs323-compilers:~/cs323/lab2$ ./wc.out lex.l
#lines #words #chars file path
40     167     1242  lex.l
cs323@deb-cs323-compilers:~/cs323/lab2$ ./wc.out lex.yy.c
#lines #words #chars file path
1776   6735   45322 lex.yy.c
```

# A Closer Look

```
1 %{
2     // just let you know you have macros!
3     // C macro tutorial in Chinese: http://c.biancheng.net/view/446.html
4     #define EXIT_OK 0
5     #define EXIT_FAIL 1
6
7     // global variables
8     int chars = 0;
9     int words = 0;
10    int lines = 0;
11 %}
12 letter [a-zA-Z]
13
14 %%%
15 {letter}+ { words++; chars+=strlen(yytext); }
16 \n { chars++; lines++; }
17 . { chars++; }
18
19 %%%
20 int main(int argc, char **argv){
21     char *file_path;
22     if(argc < 2){
23         fprintf(stderr, "Usage: %s <file_path>\n", argv[0]);
24         return EXIT_FAIL;
25     } else if(argc == 2){
26         file_path = argv[1];
27         if(!(yyin = fopen(file_path, "r"))){
28             perror(argv[1]);
29             return EXIT_FAIL;
30         }
31         yylex();
32         printf("%-8s%-8s%-8s\n", "#lines", "#words", "#chars", "file path");
33         printf("%-8d%-8d%-8d\n", lines, words, chars, file_path);
34         return EXIT_OK;
35     } else{
36         fputs("Too many arguments! Expected: 2.\n", stderr);
37         return EXIT_FAIL;
38     }
39 }
```

The structure is the same as in a Lex program:

1. Declaration
2. Translation rules
3. Auxiliary functions

# More on Flex patterns (self-study)

Flex supports a rich set of conveniences:

Character classes	<b>[0-9]</b>	This means alternation of the characters in the range listed (in this case: 0   1   2   3   4   5   6   7   8   9). More than one range may be specified, e.g. <b>[0-9A-Za-z]</b> as well as specifying individual characters, as with <b>[æiou0-9]</b> .
Character exclusion	<b>^</b>	The first character in a character class may be ^ to indicate the complement of the set of characters specified. For example, <b>[^0-9]</b> matches any non-digit character.
Arbitrary character	<b>.</b>	The period matches any single character <b>except newline</b> .
Single repetition	<b>x?</b>	0 or 1 occurrence of <b>x</b> .

# More on Flex patterns (self-study)

Non-zero repetition	<b>x+</b>	x repeated one or more times; equivalent to <b>xx*</b> .
Specified repetition	<b>x{<i>n,m</i>}</b>	x repeated between <i>n</i> and <i>m</i> times.
Beginning of line	<b>^x</b>	Match x at beginning of line only.
End of line	<b>x\$</b>	Match x at end of line only.
Context-sensitivity	<b>ab/cd</b>	Match <b>ab</b> but only when followed by <b>cd</b> . The lookahead characters are left in the input stream to be read for the next token.
Literal strings	<b>"x"</b>	This means <b>x</b> even if <b>x</b> would normally have special meaning. Thus, " <b>x*</b> " may be used to match <b>x</b> followed by an asterisk. You can turn off the special meaning of just one character by preceding it with a backslash, .e.g. \. matches exactly the period character and nothing more.
Definitions	<b>{name}</b>	Replace with the earlier defined pattern called <b>name</b> . This kind of substitution allows you to reuse pattern pieces and define more readable patterns.

<https://web.stanford.edu/class/archive/cs/cs143/cs143.1128/handouts/050%20Flex%20In%20A%20Nutshell.pdf>