



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Chapter 1: Regular Expressions & Lexical Analysis

Yepang Liu

liuyp1@sustech.edu.cn

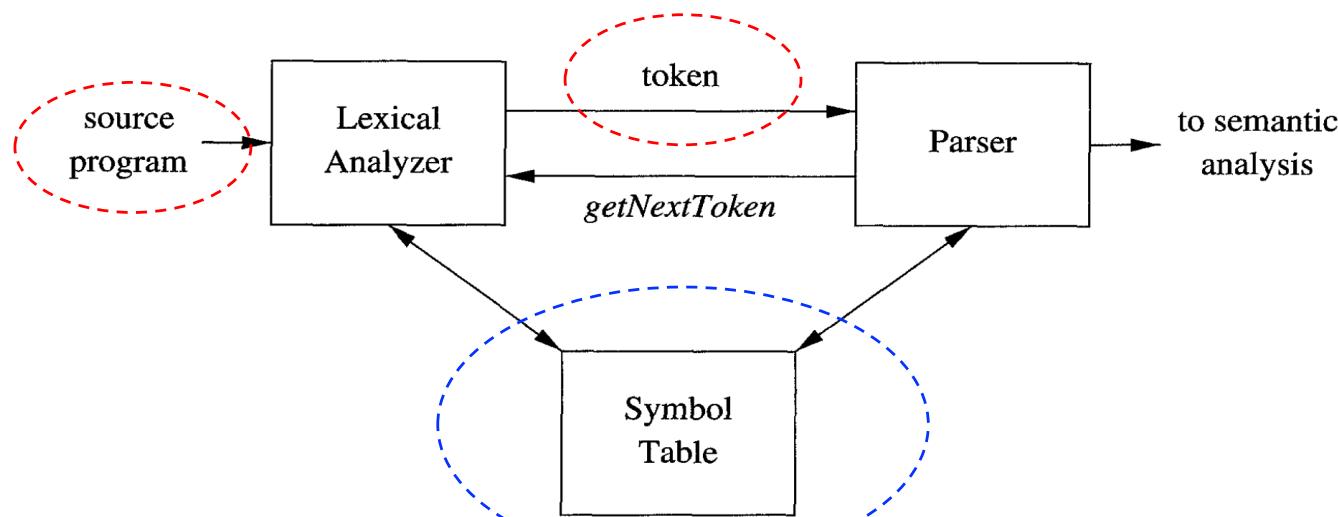
The chapter numbering in lecture notes does not follow that in the textbook.

Outline

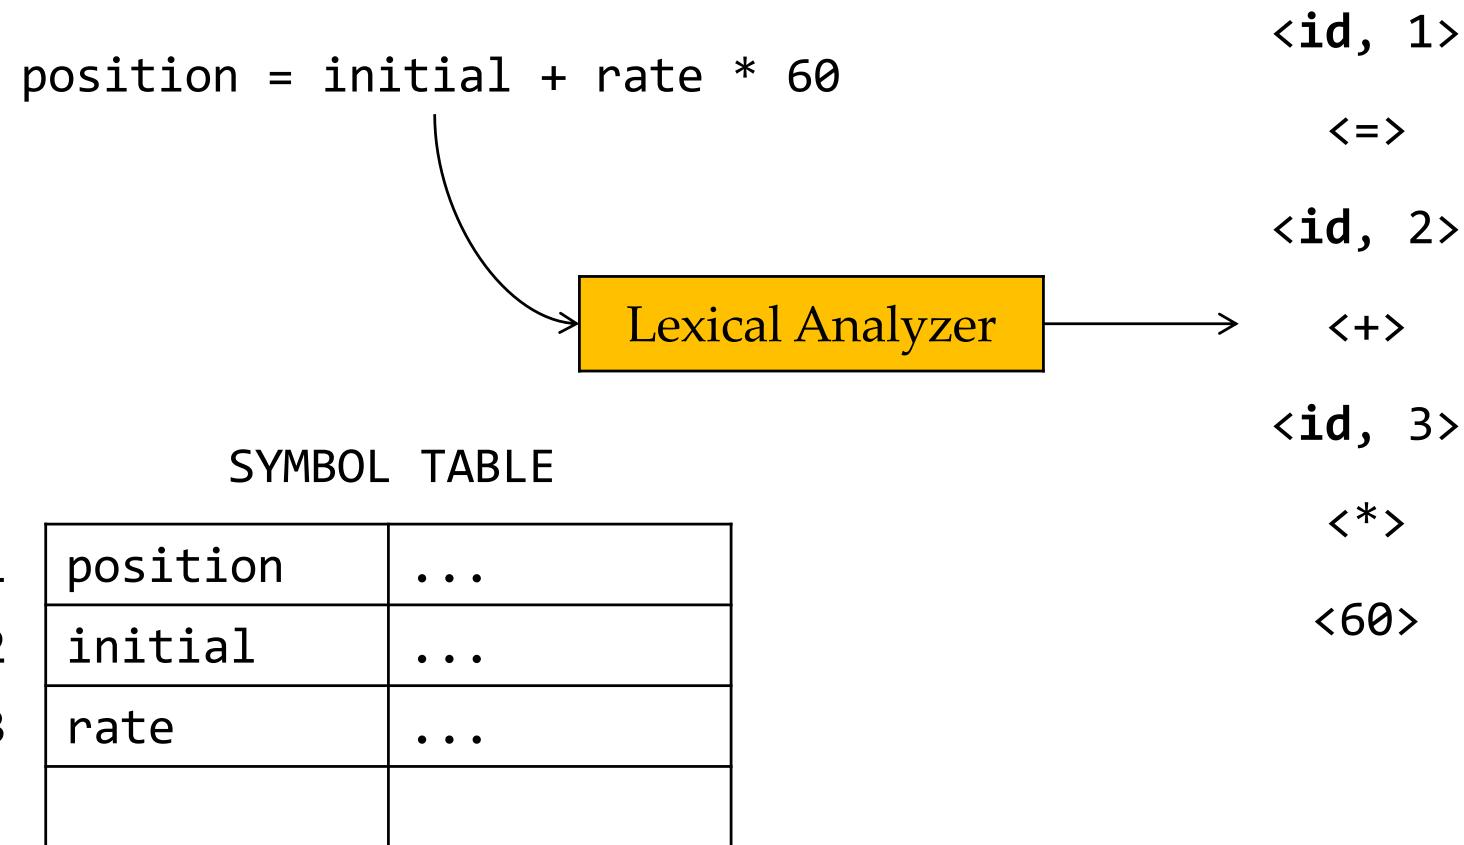
- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)

The Role of Lexical Analyzer

- Read the input characters of the source program, group them into lexemes, and produces a sequence of tokens (to be used by parser)
- Add lexemes into the symbol table when necessary



Recall the Earlier Example



Tokens, Patterns, and Lexemes

- A *lexeme* is a string of characters that is a lowest-level syntactic unit in programming languages
- A *token* is a syntactic category representing a class of lexemes. Formally, it is a pair <token name, attribute value>
 - **Token name**: an abstract symbol representing the kind of the token
 - **Attribute value** (optional) points to the symbol table
- Each token has a particular *pattern*: a description of the form that the lexemes of the token may take

Examples

Patterns

| TOKEN | INFORMAL DESCRIPTION | SAMPLE LEXEMES |
|------------|---------------------------------------|---------------------|
| if | characters i, f | if |
| else | characters e, l, s, e | else |
| comparison | < or > or <= or >= or == or != | <=, != |
| id | letter followed by letters and digits | pi, score, D2 |
| number | any numeric constant | 3.14159, 0, 6.02e23 |
| literal | anything but ", surrounded by "'s | "core dumped" |

Attributes for Tokens

- Typically, for each token, the lexical analyzer also provides additional *attribute values* to the subsequent compiler phases
 - Token names influence parsing decisions
 - Attribute values influence semantic analysis, code generation etc.
- For example, an **id** token is often associated with: (1) its lexeme, (2) data type, and (3) the location at which it is first found.
- Token attributes are stored in the **symbol table**.

A = B * 2 →

<**id**, pointer to symbol-table entry for A>
<**assign_op**>
<**id**, pointer to symbol-table entry for B>
<**mult_op**> <**number**, integer value 2>

Lexical Errors

- When the prefix of the remaining input does not match any token patterns
- Example: int **@3** = 5;

Outline

- The Role of Lexers: Recognizing Tokens
- Regular Expressions (for specifying tokens)
- Finite Automata (for recognizing patterns)

Specification of Tokens

- **Regular expression** (正则表达式, **regexp** for short) is an important notation for specifying token patterns
- Content of this part
 - Strings and Languages (串和语言)
 - Operations on Languages (语言上的运算)
 - Regular Expressions
 - Regular Definitions (正则定义)
 - Extensions of Regular Expressions

Strings and Languages

- **Alphabet (字母表)**: any finite set of symbols
 - Examples of symbols: letters, digits, and punctuations
 - Examples of alphabets: {1, 0}, ASCII, Unicode
- A **string (串)** over an alphabet is a finite sequence of symbols drawn from the alphabet
 - The length of a string s , denoted $|s|$, is the number of symbols in s (i.e., cardinality)
 - Empty string (空串): the string of length 0, ϵ

Terms (using **banana** for illustration)

- **Prefix (前缀) of string s :** any string obtained by removing 0 or more symbols from the end of s (**ban**, **banana**, ϵ)
- **Proper prefix (真前缀):** a prefix that is not ϵ and not s itself (**ban**)
- **Suffix (后缀):** any string obtained by removing 0 or more symbols from the beginning of s (**nana**, **banana**, ϵ).
- **Proper suffix (真后缀):** a suffix that is not ϵ and not equal to s itself (**nana**)

Terms Cont.

- **Substring** (子串) of s : any string obtained by removing any prefix and any suffix from s (**banana, nan, ϵ**)
- **Proper substring** (真子串): a substring that is not ϵ and not equal to s itself (**nan**)
- **Subsequence** (子序列): any string formed by removing 0 or more not necessarily consecutive symbols from s (**bnn**)



Think about this after class: How many substrings and subsequences does **banana** have?

(Two substrings are different if they have different start/end index)

String Operations (串的运算)

- **Concatenation (连接):** the concatenation of two strings x and y , denoted xy , is the string formed by appending y to x
 - $x = \text{I}, y = \heartsuit, z = \text{compilers}, xyz = \text{I}\heartsuit\text{compilers}$
- **Exponentiation (幂/指数运算):** $s^0 = \epsilon, s^1 = s, s^i = s^{i-1}s$
 - $x = 6, x^0 = \epsilon, x^1 = 6, x^3 = 666$

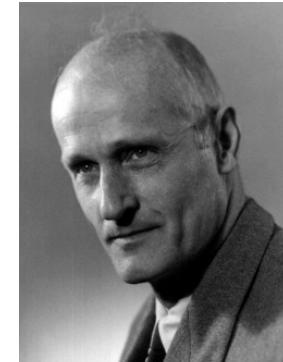
Language (语言)

- A **language** is any **countable set**¹ of strings over some fixed alphabet
 - The set containing only the empty string, that is $\{\epsilon\}$, is a language, denoted \emptyset
 - The set of all grammatically correct English sentences
 - The set of all syntactically well-formed C programs

¹In mathematics, a countable set is a set with the same cardinality (number of elements) as some subset of the set of natural numbers. A countable set is either a finite set or a countably infinite set.

Operations on Languages (语言的运算)

- Union (并), Concatenation (连接)
- Kleene Closure (Kleene闭包)
- Positive Clousre (正闭包)



Stephen C. Kleene

| OPERATION | DEFINITION AND NOTATION |
|---------------------------------|---|
| <i>Union of L and M</i> | $L \cup M = \{s \mid s \text{ is in } L \text{ or } s \text{ is in } M\}$ |
| <i>Concatenation of L and M</i> | $LM = \{st \mid s \text{ is in } L \text{ and } t \text{ is in } M\}$ |
| <i>Kleene closure of L</i> | $L^* = \bigcup_{i=0}^{\infty} L^i$ |
| <i>Positive closure of L</i> | $L^+ = \bigcup_{i=1}^{\infty} L^i$ |

The exponentiation of L can be defined using concatenation. L^n means concatenating L n times.

https://en.wikipedia.org/wiki/Stephen_Cole_Kleene

Examples

- $L = \{A, B, \dots, Z, a, b, \dots, z\}$ ----- 52 English letters
- $D = \{0, 1, \dots, 9\}$ ----- 10 digits

| | |
|-----------------|---|
| $L \cup D$ | $\{A, B, \dots, Z, a, b, \dots, z, 0, 1, \dots, 9\}$ |
| LD | the set of 520 strings of length two, each consisting of one letter followed by one digit |
| L^4 | the set of all 4-letter strings |
| L^* | the set of all strings of letters, including ϵ |
| $L(L \cup D)^*$ | ? |
| D^+ | ? |

Note: L, D might seem to be the alphabets of letters and digits. We define them to be languages: all strings happen to be of length one.

Regular Expressions - For Describing Languages/Patterns

Rules that define regexps over an alphabet Σ :

- **BASIS:** two rules form the basis:
 - ϵ is a regexp, $L(\epsilon) = \{\epsilon\}$
 - If a is a symbol in Σ , then a is a regexp, and $L(a) = \{a\}$
- **INDUCTION:** Suppose r and s are regexps denoting the languages $L(r)$ and $L(s)$
 - $(r)|(s)$ is a regexp denoting the language $L(r) \cup L(s)$
 - $(r)(s)$ is a regexp denoting the language $L(r)L(s)$
 - $(r)^*$ is a regexp denoting $(L(r))^*$
 - (r) is a regexp denoting $L(r)$, that is, additional parentheses do not change the language an expression denotes.

Regular Expressions Cont.

- Following the rules, regexps often contain **unnecessary pairs of parentheses**. We may drop some if we adopt the conventions:
 - **Precedence (优先级):** closure $*$ > concatenation > union $|$
 - **Associativity (结合性):** All three operators are left associative, meaning that operations are grouped from the left.
 - For example, $a \mid b \mid c$ would be interpreted as $(a \mid b) \mid c$
- Example: $(a) \mid ((b)^*(c))$ can be simplified as $a \mid b^*c$

Regular Expressions Examples

- Let $\Sigma = \{a, b\}$
 - $a|b$ denotes the language $\{a, b\}$
 - $(a|b)(a|b)$ denotes $\{aa, ab, ba, bb\}$
 - a^* denotes $\{\epsilon, a, aa, aaa, \dots\}$
 - $(a|b)^*$ denotes the set of all strings consisting of 0 or more a 's or b 's: $\{\epsilon, a, b, aa, ab, ba, bb, aaa, \dots\}$
 - $a|a^*b$ denotes the string a and all strings consisting of 0 or more a 's and ending in b : $\{a, b, ab, aab, aaab, \dots\}$

Regular Language (正则语言)

- A **regular language** is a language that can be defined by a regexp
- If two regexps r and s denote the same language, they are *equivalent*, written as $r = s$

$$(a|b)(a|b)$$

=

$$aa|ab|ba|bb$$

?

Algebraic Laws

SELF LEARNING
-MATERIALS

- Each law below asserts that expressions of two different forms are equivalent

| LAW | DESCRIPTION |
|----------------------------------|--|
| $r s = s r$ | is commutative |
| $r (s t) = (r s) t$ | is associative |
| $r(st) = (rs)t$ | Concatenation is associative |
| $r(s t) = rs rt; (s t)r = sr tr$ | Concatenation distributes over |
| $\epsilon r = r\epsilon = r$ | ϵ is the identity for concatenation |
| $r^* = (r \epsilon)^*$ | ϵ is guaranteed in a closure |
| $r^{**} = r^*$ | * is idempotent |

| can be viewed as + in arithmetics, concatenation can be viewed as \times , * can be viewed as the power operator.

Regular Definitions (正则定义)

- For **notational convenience**, we can give **names** (e.g., d_i below) to certain regexps and use those names in subsequent expressions

If Σ is an alphabet of basic symbols, then a **regular definition** is a sequence of definitions of the form:

$$d_1 \rightarrow r_1$$

$$d_2 \rightarrow r_2$$

...

$$d_n \rightarrow r_n$$

where:

- Each d_i is a new symbol not in Σ and not the same as the other d 's
- Each r_i is a regexp over the alphabet $\Sigma \cup \{d_1, d_2, \dots, d_{i-1}\}$

Each new symbol denotes a regular language. The second rule means that you may reuse previously-defined symbols.

Examples

- Regular definition for C identifiers

$$\begin{array}{lcl} letter_- & \rightarrow & A \mid B \mid \dots \mid Z \mid a \mid b \mid \dots \mid z \mid - \\ digit & \rightarrow & 0 \mid 1 \mid \dots \mid 9 \\ id & \rightarrow & letter_- (letter_- \mid digit)^* \end{array}$$

_hello valid?
3times valid?

- Regexp for C identifiers

$$(A|B|\dots|Z|a|b|\dots|z|_*)((A|B|\dots|Z|a|b|\dots|z|_*)|(\theta|1|\dots|9))^*$$